

Deep Learning Adventures

TensorFlow In Practice -

Presentation 5

A quick overview of Coursera's Tensorflow in Practice specialization course

Robert Kraig, David Patton, George Zoto

<https://www.meetup.com/Deep-Learning-Adventures>

In the beginning...

Meetup

The screenshot shows the homepage of the Deep Learning Adventures Meetup group. At the top, there's a photo of two men sitting on a blue couch. Below the photo, the group's name "Deep Learning Adventures" is displayed in large, bold letters. To the right of the name are links for "Start a new group", "Log in", and "Sign up". Under the group name, it says "Washington, DC", "377 members · Public group", and "Organized by George Z. and 2 others". There are "Share" buttons for Facebook, Twitter, and LinkedIn. A red "Join this group" button is prominently displayed. Below the main header, there are sections for "What we're about", "Upcoming events (2)", and "Members (377)". The "What we're about" section contains a detailed paragraph about the group's mission and history. The "Upcoming events" section lists an event with Laurence Moroney on July 3rd at 7:30 PM EDT. The "Members" section shows a grid of member profiles.

Deep Learning Adventures

Washington, DC · 377 members · Public group · Organized by George Z. and 2 others

Share: [Facebook](#) [Twitter](#) [LinkedIn](#)

[Join this group](#)

What we're about

Deep Learning Adventures is a welcoming group for anyone interested in learning more about deep learning, its foundations, its strengths and weaknesses and ever growing applications that best serve humanity and help those in need throughout the world. After participating in hundreds of meetups in the area, we have taken many lessons learned and incorporated them into this group. This group is also startup oriented in the sense that we are open minded and ready to pivot to new directions as our community and needs around the world guide us....

[Read more](#)

Upcoming events (2)

FRI, JUL 3, 7:30 PM EDT

A chat with Laurence Moroney, AI Lead at Google

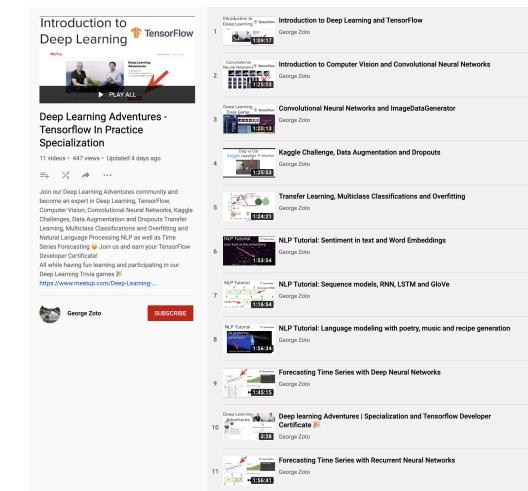
Online event

Join us for a fun conversation with Laurence Moroney, AI Lead at Google (<https://www.linkedin.com/in/laurence-moroney/>) and developer of our TensorFlow in Practice and TensorFlow: Data and Deployment Specializations! We plan to...

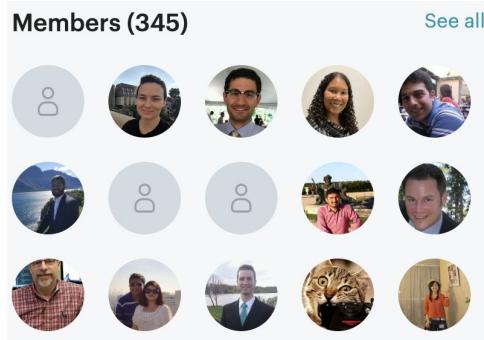
Members (377)

See all

- **Meetup Link**
<https://www.meetup.com/Deep-Learning-Adventures>
- **GitHub repository**
<https://github.com/georgezoto/Deep-Learning-Adventures>
<https://github.com/georgezoto/TensorFlow-in-Practice>
- **Join us on Slack**
https://join.slack.com/t/deeplearninga-nmk8930/shared_invite/zt-d52h9mm9-h~Q0ZXw5PXsTDzPIINivoq
- **Need a refresher or new to Deep Learning?**
- **YouTube recordings of all our Meetups 😊**
<https://bit.ly/deep-learning-tf>



Get to know our community



Hello, my name is _____ and I am a new/existing/regular member of this meetup.

I enjoy _____ and I am interested in learning more about _____

Training set 😊

Hello, my name is **George** and I am a **regular** member of this meetup. I enjoy **applying deep learning to solve interesting problems** and I am interested in learning more about **time series and forecasting**.

Not a typical Meetup... Get ready for a fun game on 7/3



To play this game

1. Use any device to open
joinmyquiz.com
2. Enter game code
775667

or share via...

START

Your Quizizz name is...

 Enter your name i 

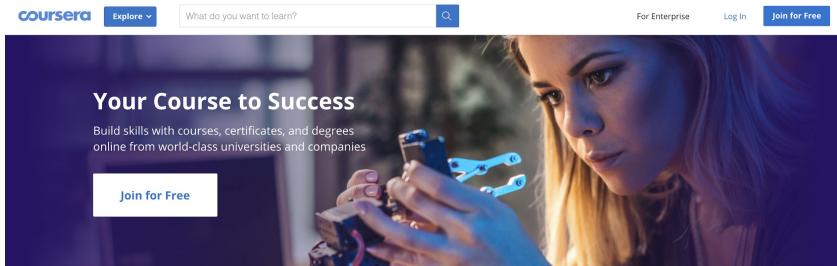
Start game

Game settings

 Music  Sound effects

 Read aloud

Attribution to Coursera and deeplearning.ai



A screenshot of the deeplearning.ai website homepage. The header features the deeplearning.ai logo and a navigation menu with links for "Courses", "Workers", "The Batch", "Events", "Forums", "Blog", and "Company". The main section has a dark background with the text "Break Into AI" and a subtext explaining that the courses will teach key concepts and applications of AI. Below this is a red button labeled "Take the Deep Learning Specialization". To the right, there's a graphic of a laptop screen showing two images: a "Content Image" of the Golden Gate Bridge and a "Generated Image" of the same bridge in a painterly style, labeled "Art Generation with Deep Learning".

Source:

<https://www.coursera.org/about/terms>
<https://www.coursera.org/>
<https://www.deeplearning.ai/>

Chapter 1 - TensorFlow in Practice Specialization

About this Specialization

199,621 recent views

Discover the tools software developers use to build scalable AI-powered algorithms in TensorFlow, a popular open-source machine learning framework.

In this four-course Specialization, you'll explore exciting opportunities for AI applications. Begin by developing an understanding of how to build and train neural networks. Improve a network's performance using convolutions as you train it to identify real-world images. You'll teach machines to understand, analyze, and respond to human speech with natural language processing systems. Learn to process text, represent sentences as vectors, and input data to a neural network. You'll even train an AI to create original poetry!

AI is already transforming industries across the world. After finishing this Specialization, you'll be able to apply your new TensorFlow skills to a wide range of problems and projects.

Looking for more advanced TensorFlow content? Check out the new [TensorFlow: Data and Deployment Specialization](#).

Chapter 1 - TensorFlow in Practice Specialization

There are 4 Courses in this Specialization

COURSE

1

Introduction to TensorFlow for Artificial Intelligence, Machine Learning, and Deep Learning

★★★★★ 4.7 6,196 ratings • 1,282 reviews

If you are a software developer who wants to build scalable AI-powered algorithms, you need to understand how to use the tools to build them. This course is part of the upcoming Machine Learning in Tensorflow Specialization and will teach you best practices for using TensorFlow, a popular open-source framework for machine learning.

[SHOW ALL](#)



6 hours to complete

A New Programming Paradigm

Welcome to this course on going from Basics to Mastery of TensorFlow. We're excited you're here! In week 1 you'll get a soft introduction to what Machine Learning and Deep Learning are, and how they offer you a new programming paradigm, giving you a new set of tools to open previously unexplored scenarios. All you need to know is some very basic

[SEE ALL](#)



4 videos (Total 16 min), 5 readings, 3 quizzes [SEE ALL](#)

COURSE

2

Convolutional Neural Networks in TensorFlow

★★★★★ 4.7 2,751 ratings • 410 reviews

If you are a software developer who wants to build scalable AI-powered algorithms, you need to understand how to use the tools to build them. This course is part of the upcoming Machine Learning in Tensorflow Specialization and will teach you best practices for using TensorFlow, a popular open-source framework for machine learning.

[SHOW ALL](#)



7 hours to complete

Introduction to Computer Vision

Welcome to week 2 of the course! In week 1 you learned all about how Machine Learning and Deep Learning is a new programming paradigm. This week you're going to take that to the next level by beginning to solve problems of computer vision with just a few lines of code!

[SEE ALL](#)



7 videos (Total 15 min), 6 readings, 3 quizzes [SEE ALL](#)

COURSE

3

Natural Language Processing in TensorFlow

★★★★★ 4.6 2,037 ratings • 277 reviews

If you are a software developer who wants to build scalable AI-powered algorithms, you need to understand how to use the tools to build them. This Specialization will teach you best practices for using TensorFlow, a popular open-source framework for machine learning.

[SHOW ALL](#)



8 hours to complete

Enhancing Vision with Convolutional Neural Networks

Welcome to week 3! In week 2 you saw a basic Neural Network for Computer Vision. It did the job nicely, but it was a little naive in its approach. This week we'll see how to make it better, as discussed by Laurence and Andrew here.

[SEE ALL](#)



6 videos (Total 19 min), 6 readings, 3 quizzes [SEE ALL](#)

COURSE

4

Sequences, Time Series and Prediction

★★★★★ 4.6 1,374 ratings • 223 reviews

If you are a software developer who wants to build scalable AI-powered algorithms, you need to understand how to use the tools to build them. This Specialization will teach you best practices for using TensorFlow, a popular open-source framework for machine learning.

[SHOW ALL](#)



9 hours to complete

Using Real-world Images

Last week you saw how to improve the results from your deep neural network using convolutions. It was a good start, but the data you used was very basic. What happens when your images are larger, or if the features aren't always in the same place? Andrew and Laurence discuss this to prepare you for what you'll learn this week: handling complex images!

[SEE ALL](#)

Source: <https://www.coursera.org/specializations/tensorflow-in-practice>

Setup



Colaboratory is a free Jupyter notebook environment that requires no setup and runs entirely in the cloud. You can write and execute code, save and share your analyses, and access powerful computing resources, all for free from your browser.

<https://colab.research.google.com>

Course 4: Sequences, Time Series and Prediction

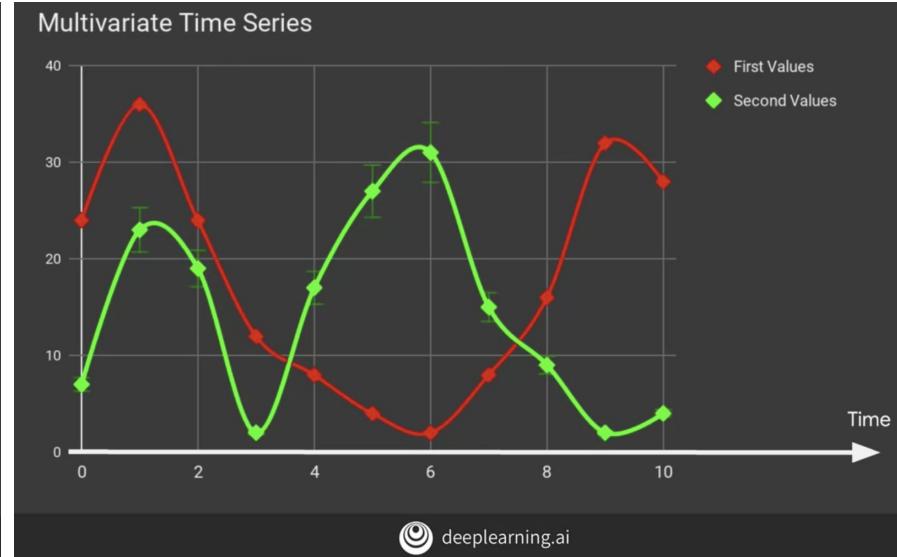
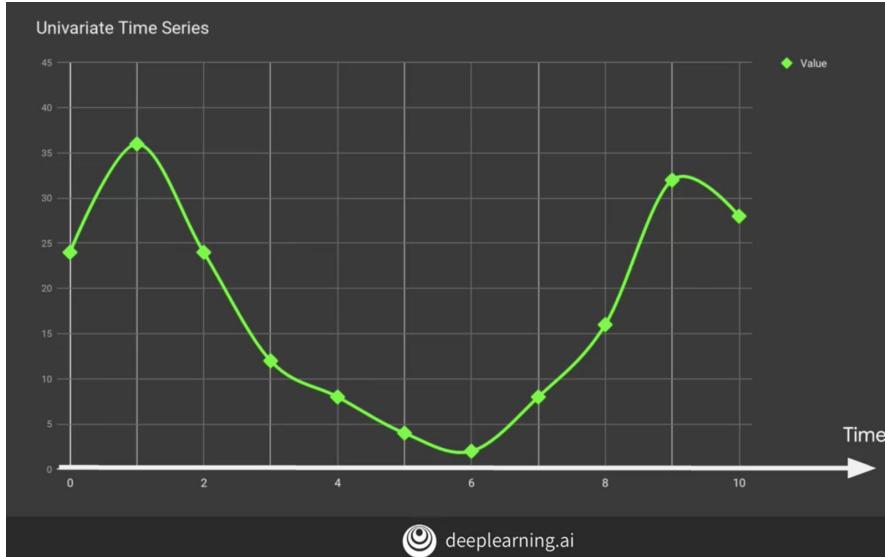
Week 1: Sequences and Prediction

Week 2: Deep Neural Networks for Time Series

Week 3: Recurrent Neural Networks for Time Series

Week 4: Real-world time series data

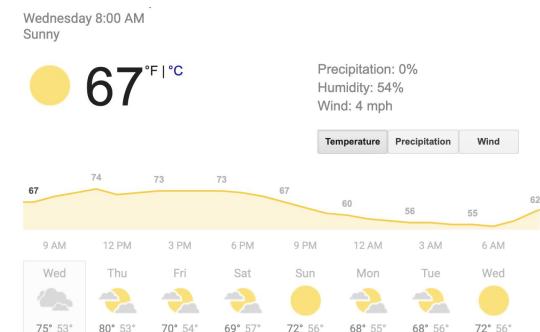
Sequences and Prediction



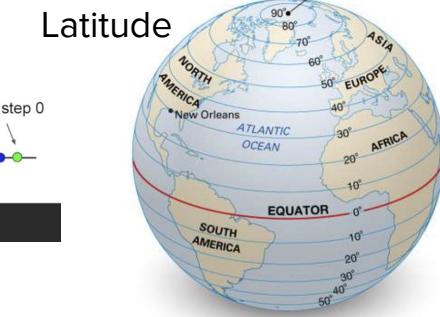
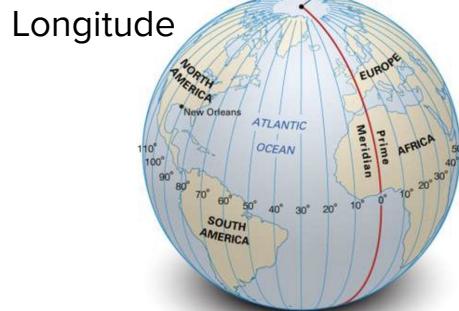
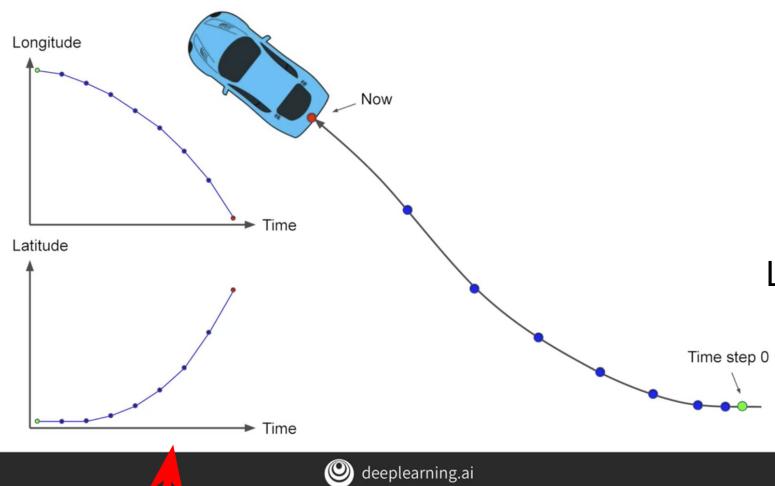
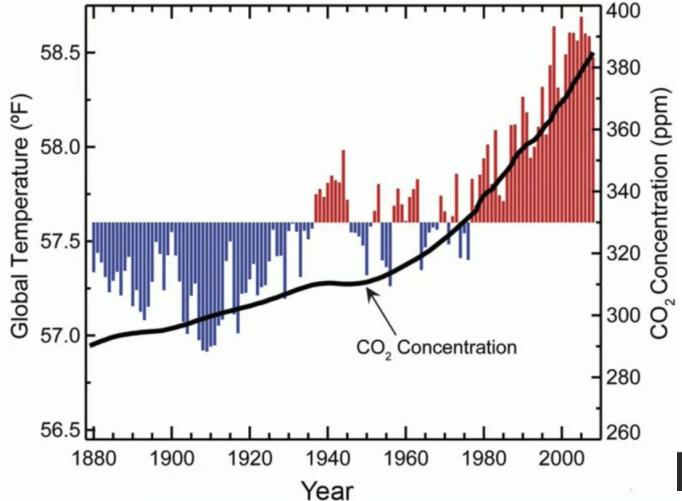
Time series: An ordered sequence of values that are usually equally spaced over time e.g. every day in the weather forecast

Single value at each time step: **univariate** time series

Multiple values at each time step: **multivariate** time series



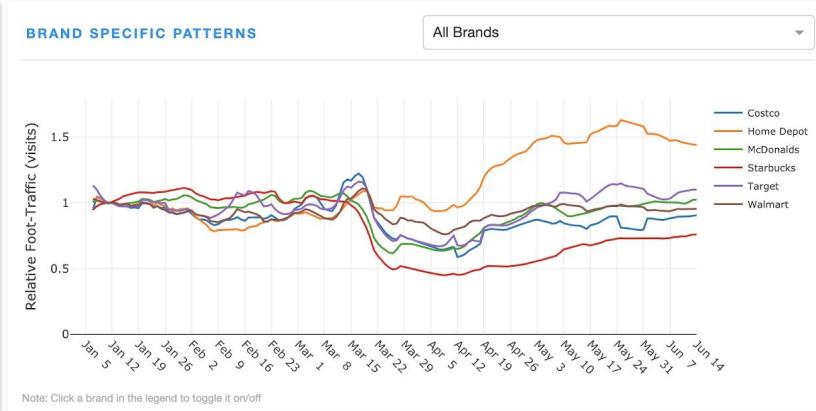
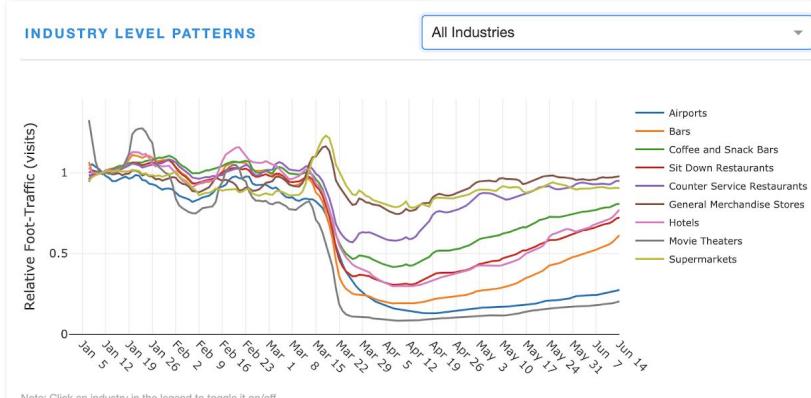
Sequences and Prediction



Movement of a body can also be plotted as a series of **univariates** or as a combined **multivariate**

Source: <https://www.coursera.org/learn/tensorflow-sequences-time-series-and-prediction>
<https://www.britannica.com/science/longitude>

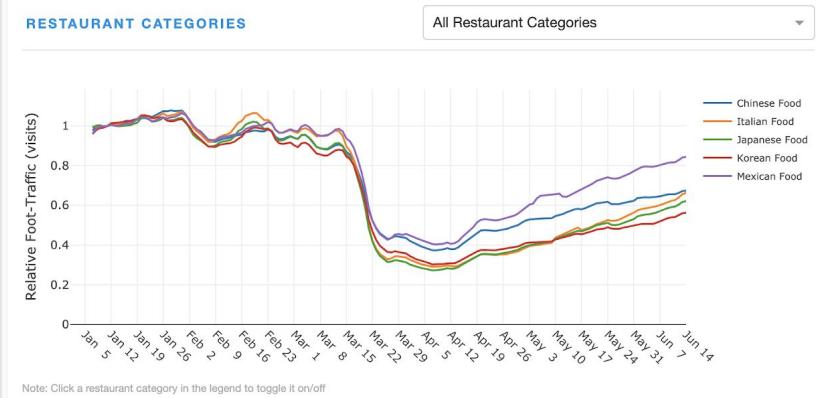
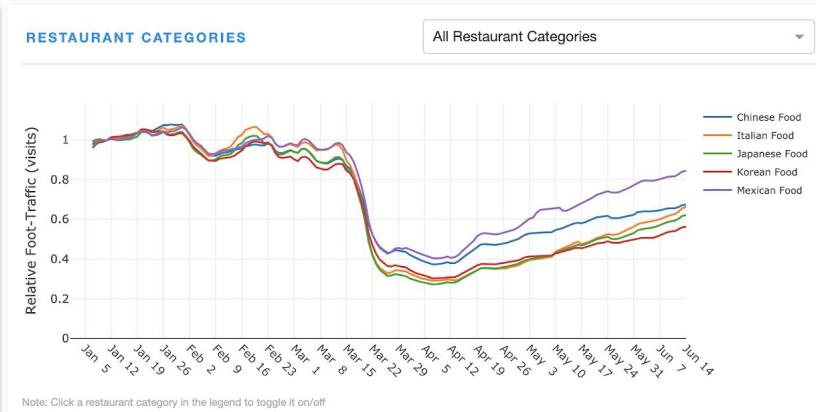
Sequences and Prediction



-8.15%

10.53%

Max



Wed, Jun 17

JN 2020

Source: <https://www.safegraph.com/dashboard/covid19-commerce-patterns>
<https://finance.google.com/finance>

Sequences and Prediction

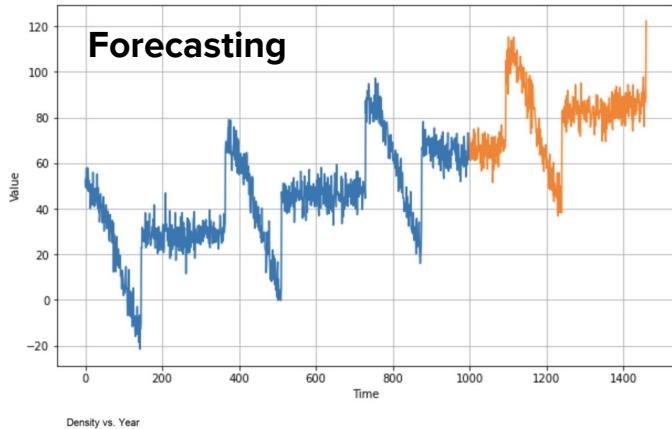


Chart Created by Imoroney@



Chart Created by Imoroney@

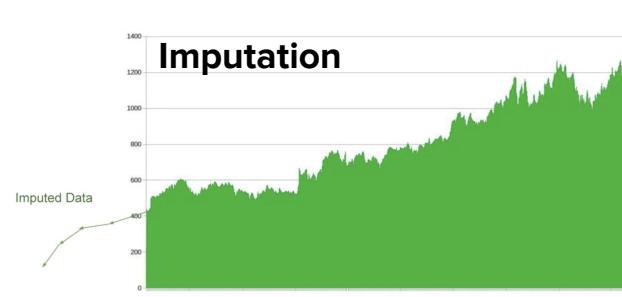
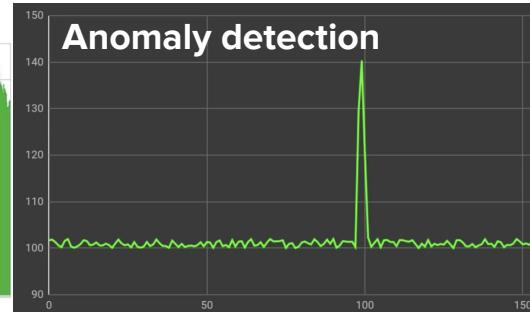
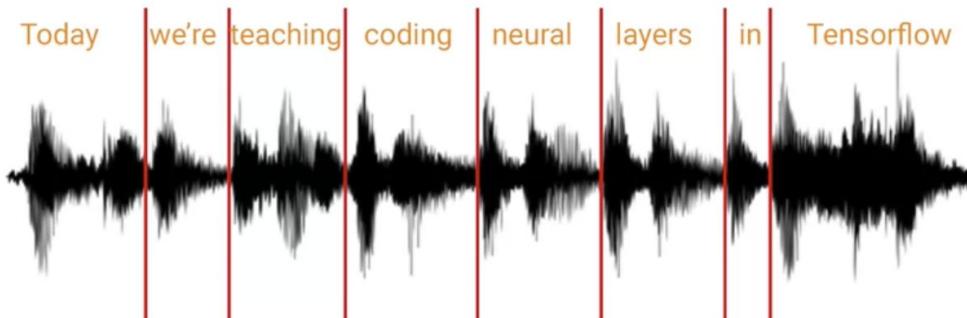


Chart Created by Imoroney@



Patterns / time series analysis



Sequences and Prediction

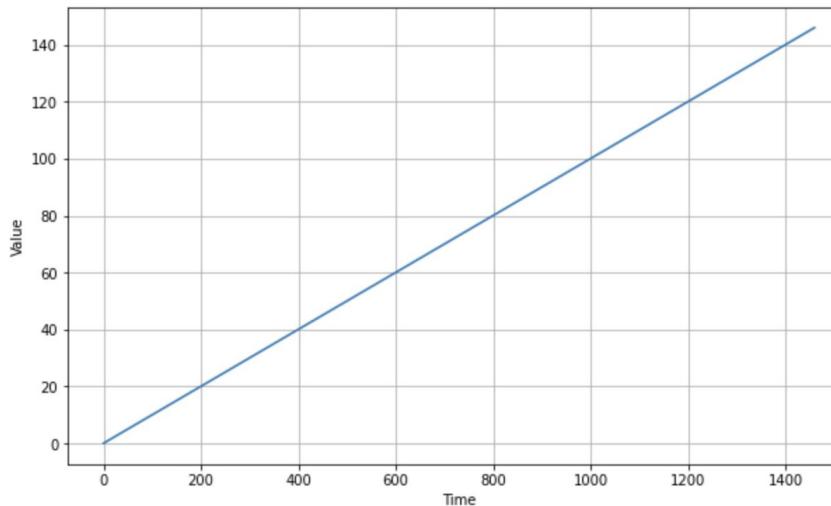
```
def trend(time, slope=0):
    return slope * time
```

```
time = np.arange(4 * 365 + 1)
print(time)

baseline = 10
series = trend(time, 0.1)
print(series)

plt.figure(figsize=(10, 6))
plot_series(time, series)
plt.show()
```

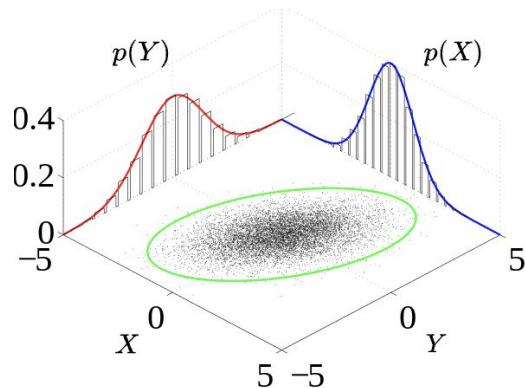
```
[ 0    1    2 ... 1458 1459 1460]
[0.000e+00 1.000e-01 2.000e-01 ... 1.458e+02 1.459e+02 1.460e+02]
```



Sequences and Prediction

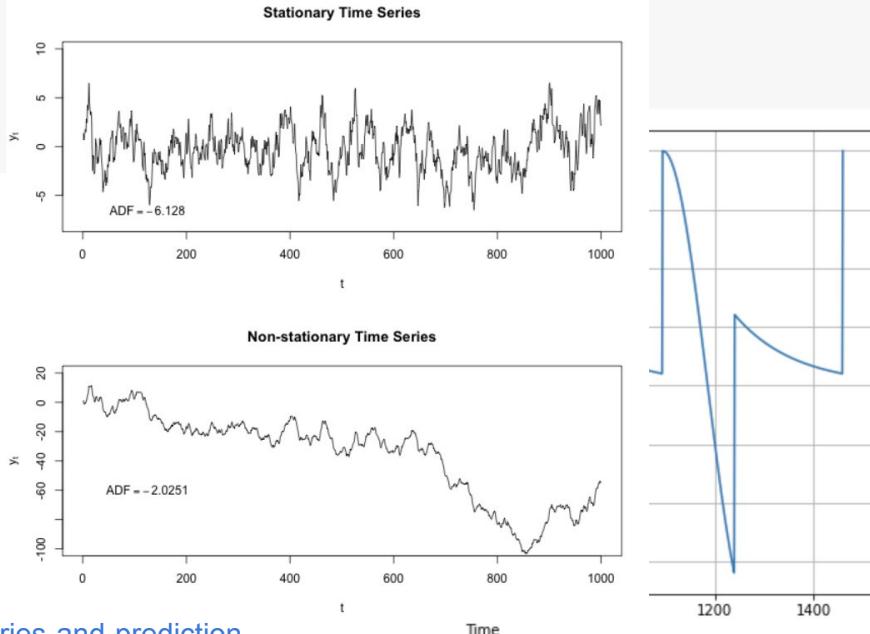
```
def seasonal_pattern(season_time):
    """Just an arbitrary pattern, you can change it if you wish"""
    return np.where(season_time < 0.4,
                   np.cos(season_time * 2 * np.pi),
                   1 / np.exp(3 * season_time))

def seasonality(time, period, amplitude=1, phase=0):
    """Repeats the same pattern at each period"""
    season_time = ((time + phase) % period) / period
    return amplitude * seasonal_pattern(season_time)
```



Stationary Time Series / Process: Its unconditional joint probability distribution does not change when shifted in time. Consequently, parameters such as mean and variance also do not change over time.

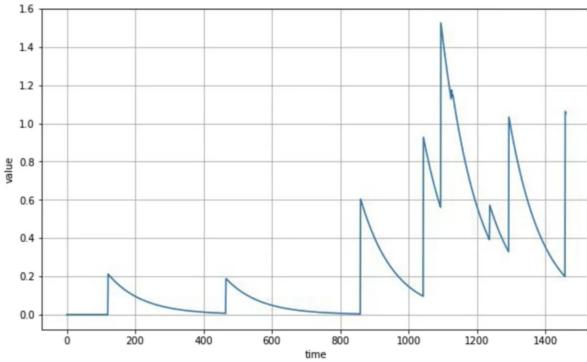
```
baseline = 10
amplitude = 40
series = seasonality(time, period=365, amplitude=amplitude)
```



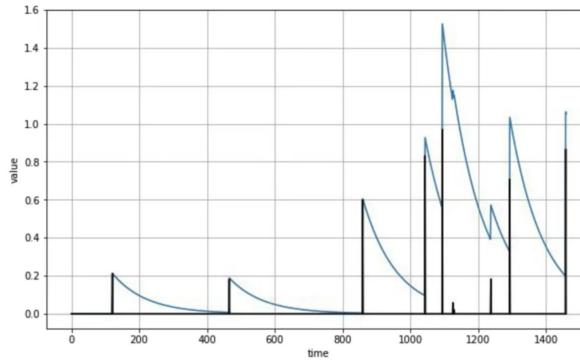
Source: <https://www.coursera.org/learn/tensorflow-sequences-time-series-and-prediction>
https://en.wikipedia.org/wiki/Stationary_process

Sequences and Prediction

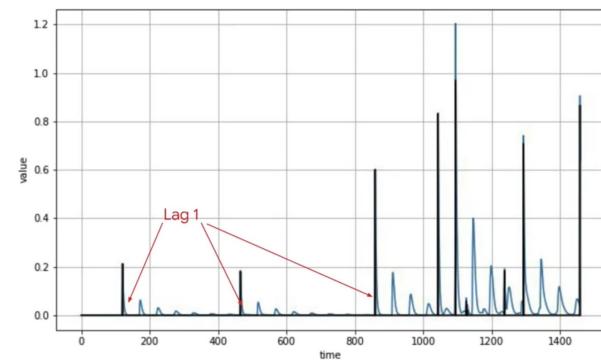
Autocorrelation



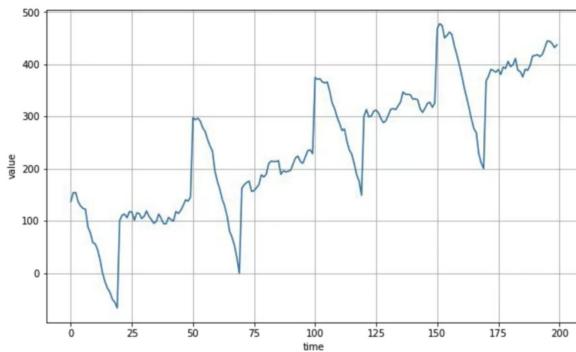
$$v(t) = 0.99 \times v(t-1) + \text{occasional spike}$$



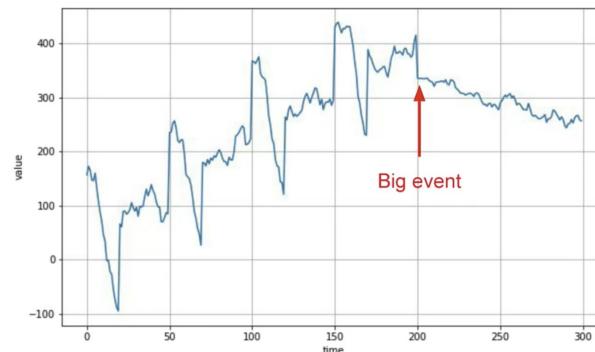
$$v(t) = 0.7 \times v(t-1) + 0.2 \times v(t-50) + \text{occasional spike}$$



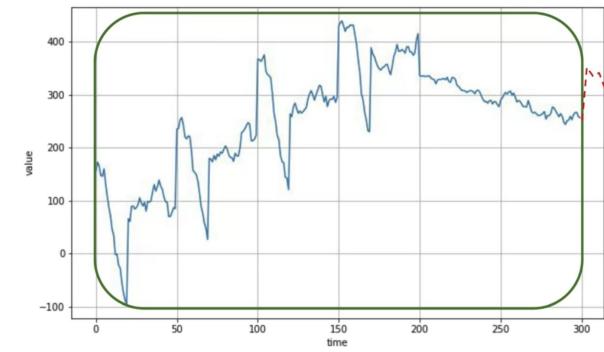
Trend + Seasonality + Autocorrelation + Noise



Non-Stationary Time Series



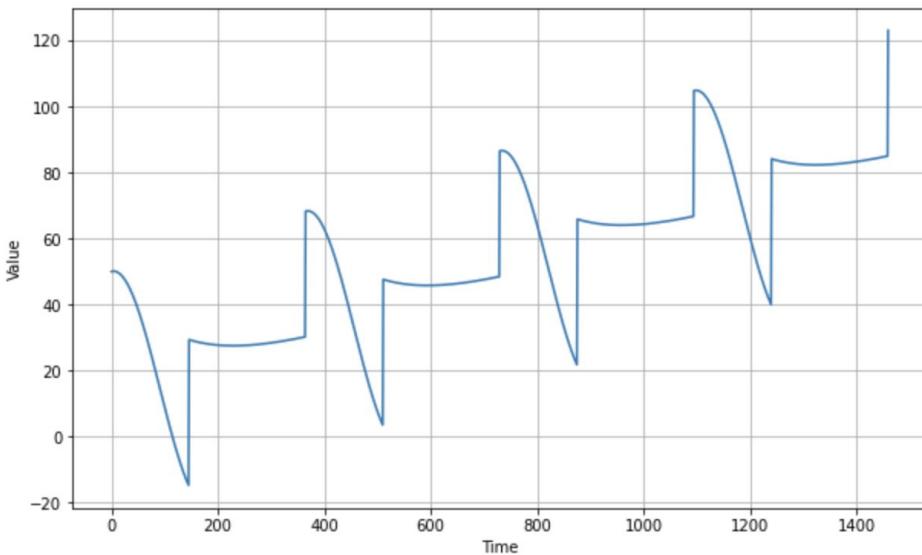
Non-Stationary Time Series



Sequences and Prediction

```
slope = 0.05
series = baseline + trend(time, slope) + seasonality(time, period=365, amplitude=amplitude)

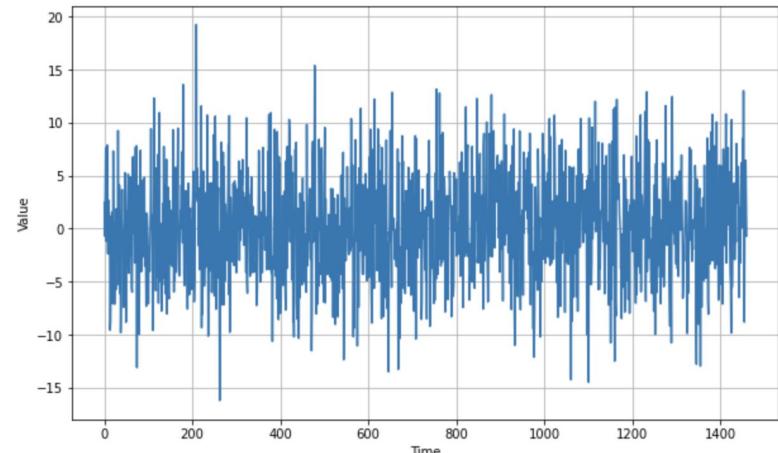
plt.figure(figsize=(10, 6))
plot_series(time, series)
plt.show()
```



```
def white_noise(time, noise_level=1, seed=None):
    rnd = np.random.RandomState(seed)
    return rnd.randn(len(time)) * noise_level

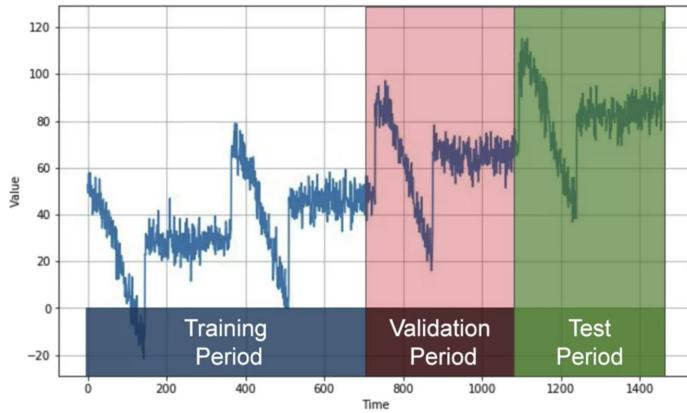
noise_level = 5
noise = white_noise(time, noise_level, seed=42)

plt.figure(figsize=(10, 6))
plot_series(time, noise)
plt.show()
```

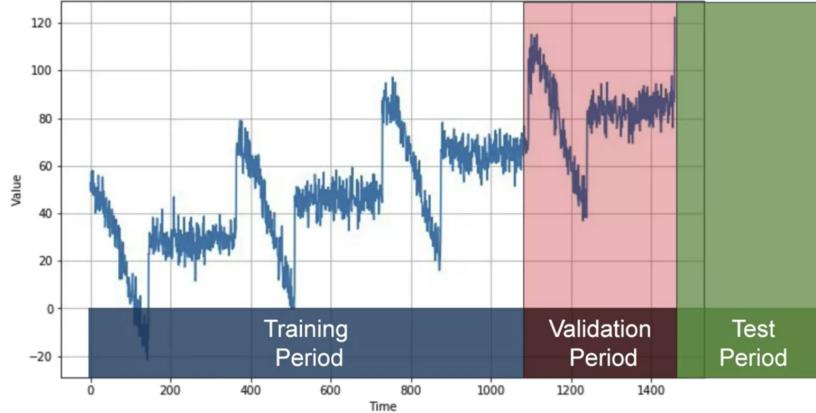
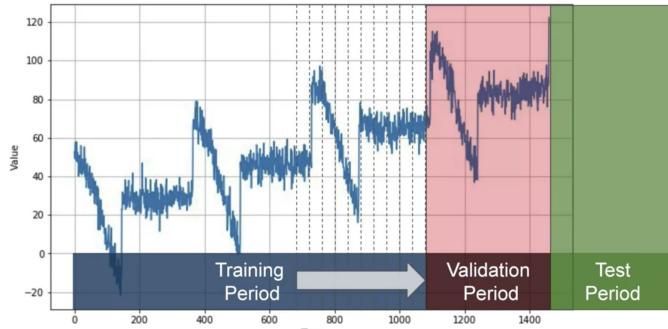


Sequences and Prediction

Fixed Partitioning



Roll-Forward Partitioning



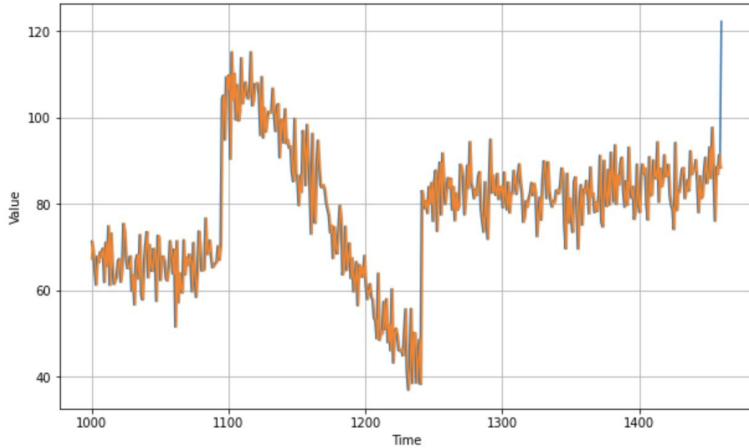
Metrics

```
errors = forecasts - actual  
  
mse = np.square(errors).mean()  
  
rmse = np.sqrt(mse)  
  
mae = np.abs(errors).mean()  
  
mape = np.abs(errors / x_valid).mean()
```

Sequences and Prediction

```
# Create the series
series = baseline + trend(time, slope) + seasonality(time, period=365, amplitude=amplitude)
# Update with noise
series += noise(time, noise_level, seed=42)
```

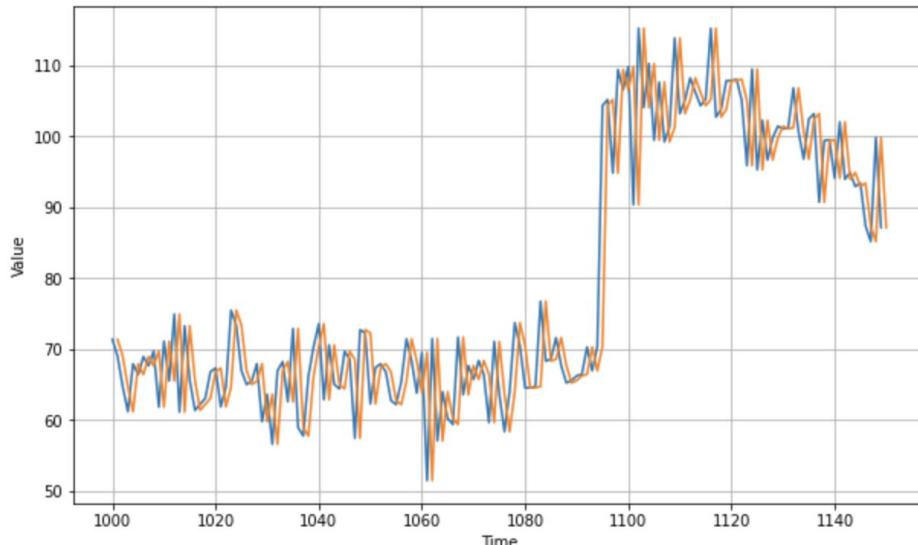
```
naive_forecast = series[split_time - 1:-1]
```



```
print(keras.metrics.mean_squared_error(x_valid, naive_forecast).numpy())
print(keras.metrics.mean_absolute_error(x_valid, naive_forecast).numpy())
```

61.827538
5.9379086

```
split_time = 1000
time_train = time[:split_time]
x_train = series[:split_time]
time_valid = time[split_time:]
x_valid = series[split_time:]
```

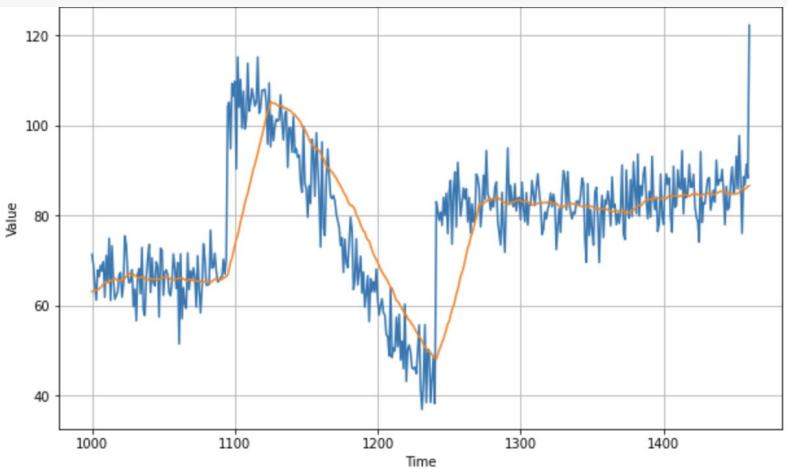


Source: <https://www.coursera.org/learn/tensorflow-sequences-time-series-and-prediction>

Sequences and Prediction

```
def moving_average_forecast(series, window_size):
    """Forecasts the mean of the last few values.
    If window_size=1, then this is equivalent to naive forecast"""
    forecast = []
    for time in range(len(series) - window_size):
        forecast.append(series[time:time + window_size].mean())
    return np.array(forecast)

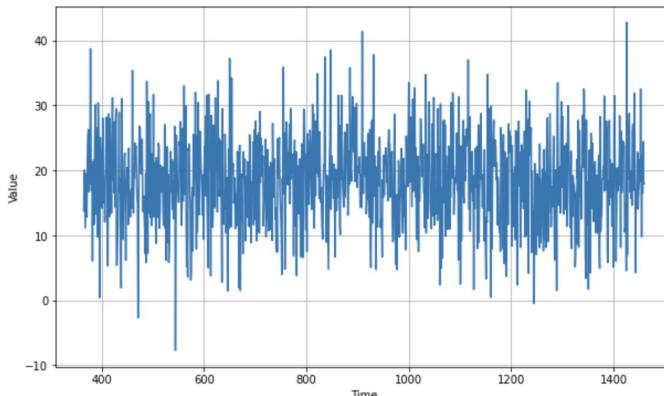
moving_avg = moving_average_forecast(series, 30)[split_time - 30:]
```



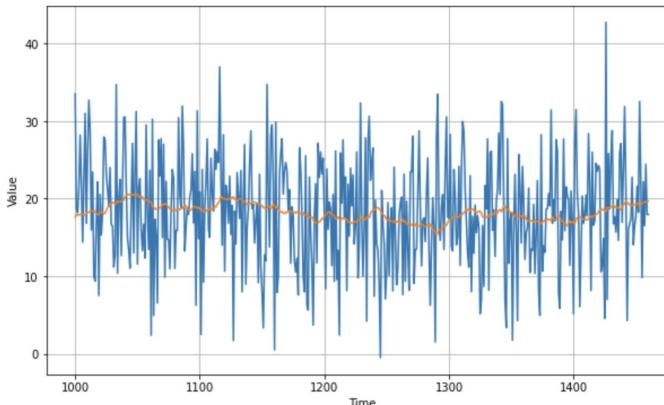
```
print(keras.metrics.mean_squared_error(x_valid, moving_avg).numpy())
print(keras.metrics.mean_absolute_error(x_valid, moving_avg).numpy())
```

106.674576
7.142419

```
diff_series = (series[365:] - series[:-365])
diff_time = time[365:]
```



```
diff_moving_avg = moving_average_forecast(diff_series, 50)[split_time - 365 - 50:]
```

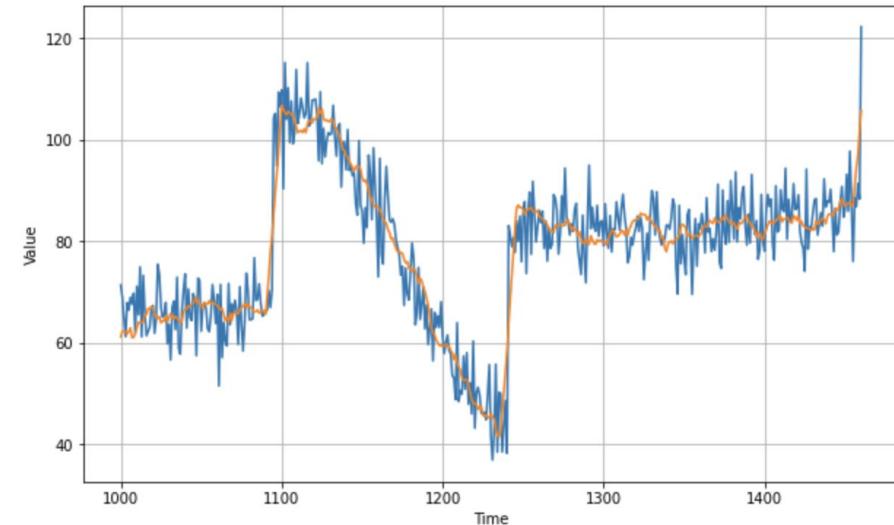
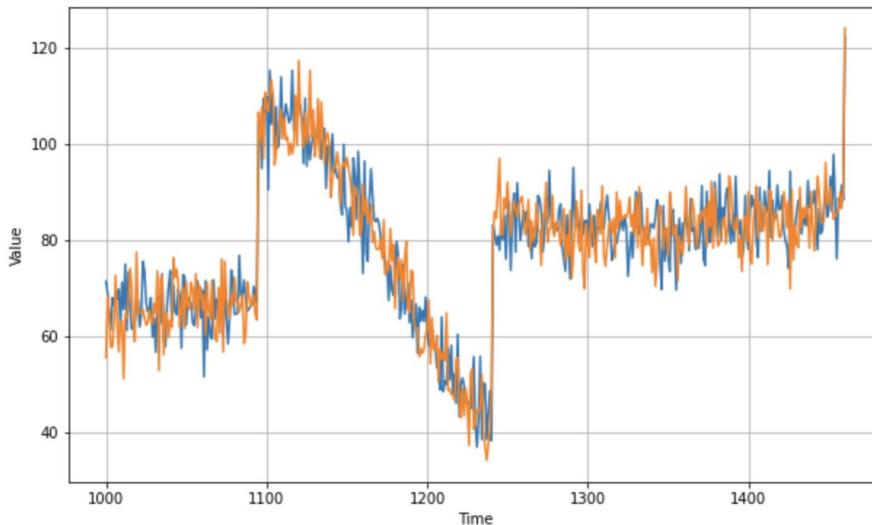


Source: <https://www.coursera.org/learn/tensorflow-sequences-time-series-and-prediction>

Sequences and Prediction

```
diff_moving_avg_plus_past = series[split_time - 365:-365] + diff_moving_avg
```

```
diff_moving_avg_plus_smooth_past = moving_average_forecast(series[split_time - 370:-360], 10) + diff_moving_avg
```



```
print(keras.metrics.mean_squared_error(x_valid, diff_moving_avg_pl)) print(keras.metrics.mean_squared_error(x_valid, diff_moving_avg_plus_smooth_past).numpy())
print(keras.metrics.mean_absolute_error(x_valid, diff_moving_avg_p)) print(keras.metrics.mean_absolute_error(x_valid, diff_moving_avg_plus_smooth_past).numpy())
```

52.973656
5.8393106

33.45226
4.569442

Source: <https://www.coursera.org/learn/tensorflow-sequences-time-series-and-prediction>

Course 4: Sequences, Time Series and Prediction

Week 1: Sequences and Prediction

Week 2: Deep Neural Networks for Time Series

Week 3: Recurrent Neural Networks for Time Series

Week 4: Real-world time series data

Deep Neural Networks for Time Series

```
dataset = tf.data.Dataset.range(10)
for val in dataset:
    print(val.numpy())
0
1
2
3
4
5
6
7
8
9
```

```
dataset = tf.data.Dataset.range(10)
dataset = dataset.window(5, shift=1)
for window_dataset in dataset:
    for val in window_dataset:
        print(val.numpy(), end=" ")
    print()
```

```
0 1 2 3 4
1 2 3 4 5
2 3 4 5 6
3 4 5 6 7
4 5 6 7 8
5 6 7 8 9
6 7 8 9
7 8 9
8 9
9
```

```
dataset = tf.data.Dataset.range(10)
dataset = dataset.window(5, shift=1, drop_remainder=True)
for window_dataset in dataset:
    for val in window_dataset:
        print(val.numpy(), end=" ")
    print()
```

```
0 1 2 3 4
1 2 3 4 5
2 3 4 5 6
3 4 5 6 7
4 5 6 7 8
5 6 7 8 9
```

```
dataset = tf.data.Dataset.range(10)
dataset = dataset.window(5, shift=1, drop_remainder=True)
dataset = dataset.flat_map(lambda window: window.batch(5))
for window in dataset:
    print(window.numpy())
```

```
[0 1 2 3 4]
[1 2 3 4 5]
[2 3 4 5 6]
[3 4 5 6 7]
[4 5 6 7 8]
[5 6 7 8 9]
```

```
dataset = tf.data.Dataset.range(10)
dataset = dataset.window(5, shift=1, drop_remainder=True)
dataset = dataset.flat_map(lambda window: window.batch(5))
dataset = dataset.map(lambda window: (window[:-1], window[-1:]))
for x,y in dataset:
    print(x.numpy(), y.numpy())

```

```
[0 1 2 3] [4]
[1 2 3 4] [5]
[2 3 4 5] [6]
[3 4 5 6] [7]
[4 5 6 7] [8]
[5 6 7 8] [9]
```

```
dataset = tf.data.Dataset.range(10)
dataset = dataset.window(5, shift=1, drop_remainder=True)
dataset = dataset.flat_map(lambda window: window.batch(5))
dataset = dataset.map(lambda window: (window[:-1], window[-1:]))
dataset = dataset.shuffle(buffer_size=10)
for x,y in dataset:
    print(x.numpy(), y.numpy())

```

```
[0 1 2 3] [4]
[2 3 4 5] [6]
[5 6 7 8] [9]
[1 2 3 4] [5]
[4 5 6 7] [8]
[3 4 5 6] [7]
```

```
dataset = tf.data.Dataset.range(10)
dataset = dataset.window(5, shift=1, drop_remainder=True)
dataset = dataset.flat_map(lambda window: window.batch(5))
dataset = dataset.map(lambda window: (window[:-1], window[-1:]))
dataset = dataset.shuffle(buffer_size=10)
dataset = dataset.batch(2).prefetch(1)
for x,y in dataset:
    print("x = ", x.numpy())
    print("y = ", y.numpy())

```

```
x = [[0 1 2 3]
      [5 6 7 8]]
y = [[4]
     [9]]
x = [[1 2 3 4]
      [2 3 4 5]]
y = [[5]
     [6]]
x = [[3 4 5 6]
      [4 5 6 7]]
y = [[7]
     [8]]
```

Deep Neural Networks for Time Series

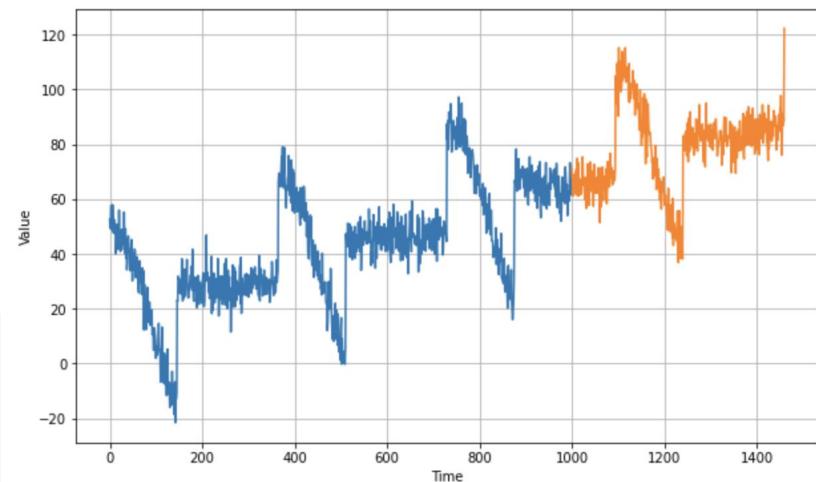
```
# Create the series
series = baseline + trend(time, slope) + seasonality(time, period=365, amplitude=amplitude)
# Update with noise
series += noise(time, noise_level, seed=42)

split_time = 1000
time_train = time[:split_time]
x_train = series[:split_time]
time_valid = time[split_time:]
x_valid = series[split_time:]

window_size = 20
batch_size = 32
shuffle_buffer_size = 1000

def windowed_dataset(series, window_size, batch_size, shuffle_buffer):
    dataset = tf.data.Dataset.from_tensor_slices(series)
    dataset = dataset.window(window_size + 1, shift=1, drop_remainder=True)
    dataset = dataset.flat_map(lambda window: window.batch(window_size + 1))
    dataset = dataset.shuffle(shuffle_buffer).map(lambda window: (window[:-1], window[-1]))
    dataset = dataset.batch(batch_size).prefetch(1)
    return dataset
```

```
time = np.arange(4 * 365 + 1, dtype="float32")
baseline = 10
series = trend(time, 0.1)
baseline = 10
amplitude = 40
slope = 0.05
noise_level = 5
```



```
dataset = windowed_dataset(x_train, window_size, batch_size, shuffle_buffer_size)
print(dataset)
10 = tf.keras.layers.Dense(1, input_shape=[window_size])
model = tf.keras.models.Sequential([10])

model.compile(loss="mse", optimizer=tf.keras.optimizers.SGD(lr=1e-6, momentum=0.9), metrics=['accuracy'])
model.fit(dataset, epochs=50, verbose=1)
```

Deep Neural Networks for Time Series

```
Epoch 40/50
31/31 [=====] - 0s 2ms/step - loss: 54.7502 - accuracy: 0.0000e+00
Epoch 41/50
31/31 [=====] - 0s 2ms/step - loss: 52.2573 - accuracy: 0.0000e+00
Epoch 42/50
31/31 [=====] - 0s 2ms/step - loss: 52.1144 - accuracy: 0.0000e+00
Epoch 43/50
31/31 [=====] - 0s 3ms/step - loss: 51.1316 - accuracy: 0.0000e+00
Epoch 44/50
31/31 [=====] - 0s 3ms/step - loss: 51.8674 - accuracy: 0.0000e+00
Epoch 45/50
31/31 [=====] - 0s 2ms/step - loss: 50.2164 - accuracy: 0.0000e+00
Epoch 46/50
31/31 [=====] - 0s 2ms/step - loss: 49.6189 - accuracy: 0.0000e+00
Epoch 47/50
31/31 [=====] - 0s 2ms/step - loss: 50.3166 - accuracy: 0.0000e+00
Epoch 48/50
31/31 [=====] - 0s 2ms/step - loss: 49.0882 - accuracy: 0.0000e+00
Epoch 49/50
31/31 [=====] - 0s 2ms/step - loss: 49.3388 - accuracy: 0.0000e+00
Epoch 50/50
31/31 [=====] - 0s 2ms/step - loss: 50.4229 - accuracy: 0.0000e+00

forecast = []

for time in range(len(series) - window_size):
    forecast.append(model.predict(series[time:time + window_size][np.newaxis]))

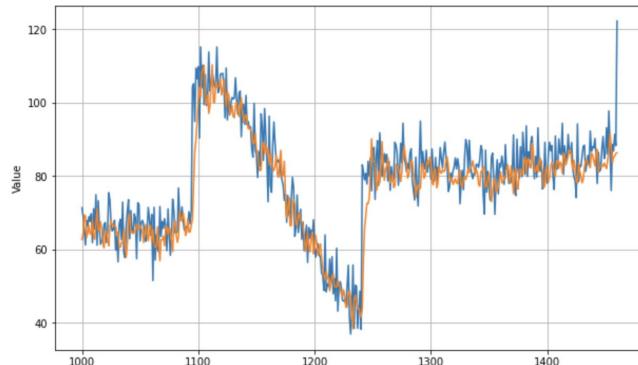
forecast = forecast[split_time - window_size:]
results = np.array(forecast)[:, 0, 0]

plt.figure(figsize=(10, 6))

plot_series(time_valid, x_valid)
plot_series(time_valid, results)
```

```
print("Layer weights {}".format(10.get_weights()))

Layer weights [array([[ 0.04292529],
       [-0.06003034],
       [-0.12568147],
       [ 0.12413348],
       [ 0.13928801],
       [-0.04836606],
       [-0.05267965],
       [ 0.05041023],
       [-0.05249092],
       [-0.17136675],
       [-0.03027576],
       [ 0.01343512],
       [ 0.09165959],
       [ 0.08489882],
       [ 0.13626154],
       [ 0.08202708],
       [-0.00438521],
       [ 0.12374628],
       [ 0.29035616],
       [ 0.34197554]], dtype=float32), array([0.01085013], dtype=float32)]
```

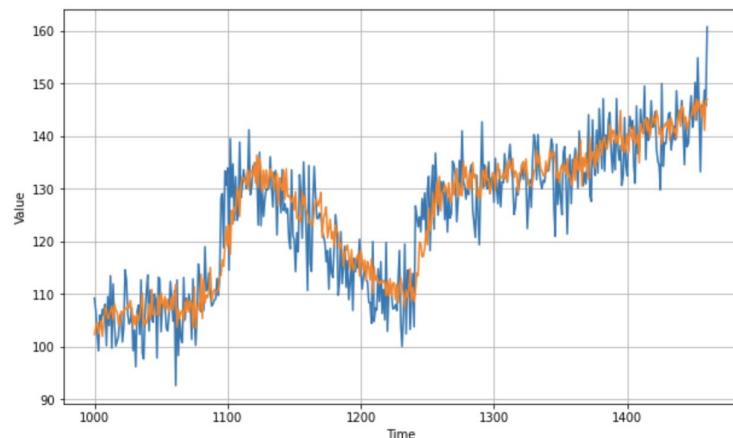
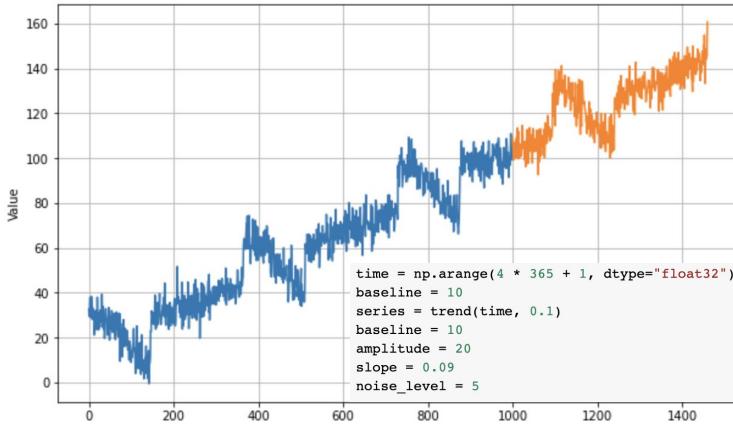


```
tf.keras.metrics.mean_absolute_error(x_valid, results).numpy()
```

5.2385697

Source: <https://www.coursera.org/learn/tensorflow-sequences-time-series-and-prediction>

Deep Neural Networks for Time Series



```
dataset = windowed_dataset(x_train, window_size, batch_size, shuffle_buffer_size)

model = tf.keras.models.Sequential([
    tf.keras.layers.Dense(10, input_shape=[window_size], activation="relu"),
    tf.keras.layers.Dense(10, activation="relu"),
    tf.keras.layers.Dense(1)
])

model.compile(loss="mse", optimizer=tf.keras.optimizers.SGD(lr=1e-6, momentum=0.9))
model.fit(dataset, epochs=my_epochs, verbose=1)
```

```
Epoch 97/100
31/31 [=====] - 0s 2ms/step - loss: 38.6709
Epoch 98/100
31/31 [=====] - 0s 2ms/step - loss: 38.8540
Epoch 99/100
31/31 [=====] - 0s 2ms/step - loss: 39.5055
Epoch 100/100
31/31 [=====] - 0s 2ms/step - loss: 38.8792
```

```
tf.keras.metrics.mean_absolute_error(x_valid, results).numpy()
```

4.678847

Deep Neural Networks for Time Series

```
dataset = windowed_dataset(x_train, window_size, batch_size, shuffle_buffer_size)

model = tf.keras.models.Sequential([
    tf.keras.layers.Dense(10, input_shape=[window_size], activation="relu"),
    tf.keras.layers.Dense(10, activation="relu"),
    tf.keras.layers.Dense(1)
])

lr_schedule = tf.keras.callbacks.LearningRateScheduler(lambda epoch: 1e-8 * 10**((epoch / 20)))

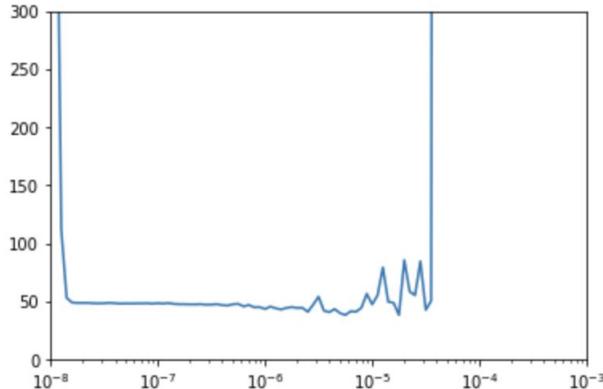
optimizer = tf.keras.optimizers.SGD(lr=1e-8, momentum=0.9)
model.compile(loss="mse", optimizer=optimizer)

history = model.fit(dataset, epochs=my_epochs, callbacks=[lr_schedule], verbose=1)
```

```
Epoch 97/100
31/31 [=====] - 0s 2ms/step - loss: 742.2804 - lr: 6.3096e-04
Epoch 98/100
31/31 [=====] - 0s 2ms/step - loss: 725.2363 - lr: 7.0795e-04
Epoch 99/100
31/31 [=====] - 0s 2ms/step - loss: 727.0015 - lr: 7.9433e-04
Epoch 100/100
31/31 [=====] - 0s 3ms/step - loss: 739.7311 - lr: 8.9125e-04
```

```
lrs = 1e-8 * (10 ** (np.arange(my_epochs) / 20))
plt.semilogx(lrs, history.history["loss"])
plt.axis([1e-8, 1e-3, 0, 300])
```

```
(1e-08, 0.001, 0.0, 300.0)
```



Deep Neural Networks for Time Series

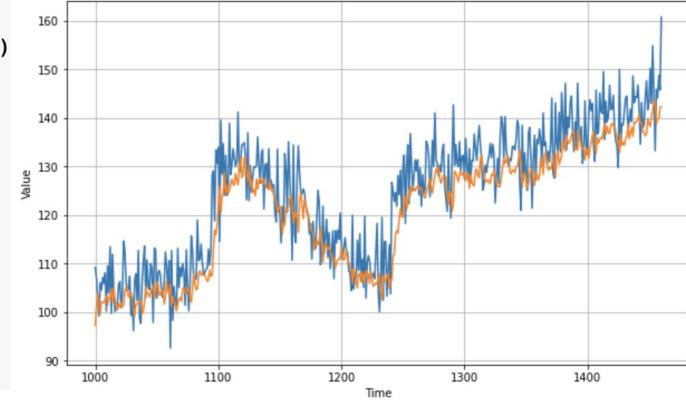
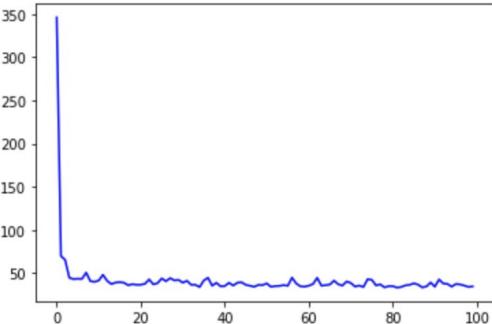
```
window_size = 30
dataset = windowed_dataset(x_train, window_size, batch_size, shuffle_buffer_size)

model = tf.keras.models.Sequential([
    tf.keras.layers.Dense(10, activation="relu", input_shape=[window_size]),
    tf.keras.layers.Dense(10, activation="relu"),
    tf.keras.layers.Dense(1)
])

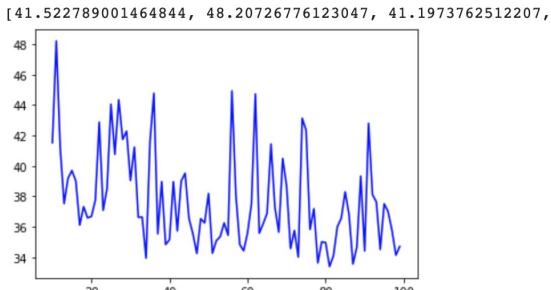
optimizer = tf.keras.optimizers.SGD(lr=8e-6, momentum=0.9)
model.compile(loss="mse", optimizer=optimizer)
history = model.fit(dataset, epochs=my_epochs, verbose=0)

loss = history.history['loss']
epochs = range(len(loss))
plt.plot(epochs, loss, 'b', label='Training Loss')
plt.show()

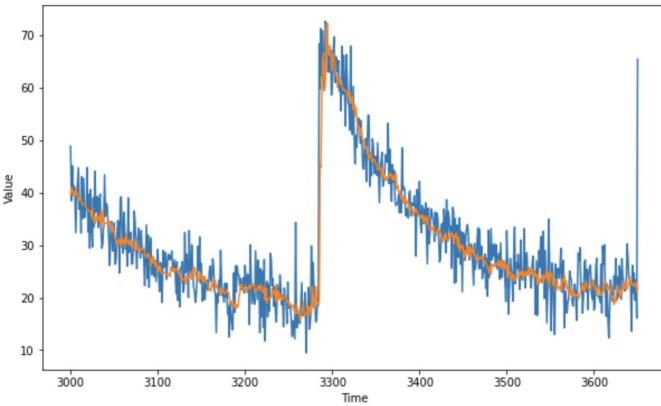
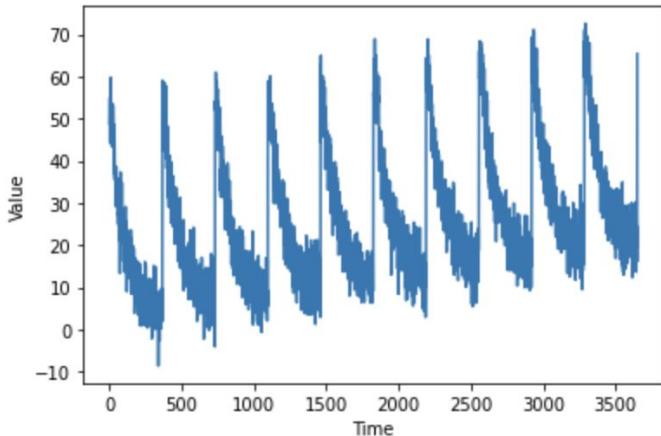
# Plot all but the first 10
loss = history.history['loss']
epochs = range(10, len(loss))
plot_loss = loss[10:]
print(plot_loss)
plt.plot(epochs, plot_loss, 'b', label='Training Loss')
plt.show()
```



```
tf.keras.metrics.mean_absolute_error(x_valid, results).numpy()
5.301772
```



Deep Neural Networks for Time Series



```
dataset = windowed_dataset(x_train, window_size, batch_size, shuffle_buffer_size)

model = tf.keras.models.Sequential([
    tf.keras.layers.Dense(10, input_shape=[window_size], activation='relu'),
    tf.keras.layers.Dense(10, activation="relu"),
    tf.keras.layers.Dense(1)
])

model.compile(loss='mse', optimizer=tf.keras.optimizers.SGD(learning_rate=1e-6, momentum=0.9))
model.fit(dataset, epochs=my_epochs, verbose=1)
```

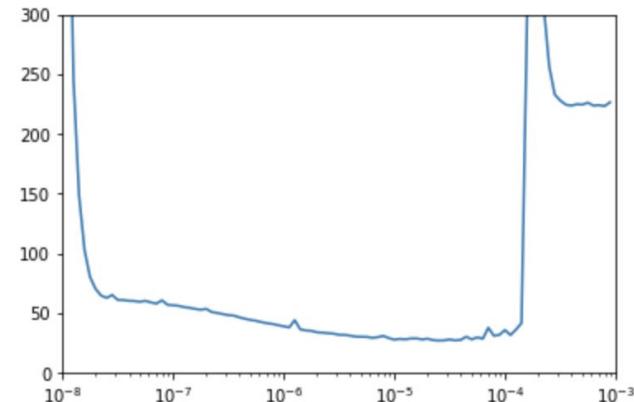
```
time = np.arange(10 * 365 + 1,
series = trend(time, 0.1)
baseline = 0
amplitude = 50
slope = 0.005
noise_level = 4
```

Total params: 331
Trainable params: 331
Non-trainable params: 0

```
tf.keras.metrics.mean_absolute_error(
# EXPECTED OUTPUT
# A Value less than 3
3.574421
```

```
lrs = 1e-8 * (10 ** (np.arange(my_epochs) / 20))
plt.semilogx(lrs, history.history["loss"])
plt.axis([1e-8, 1e-3, 0, 300])
```

(1e-08, 0.001, 0.0, 300.0)

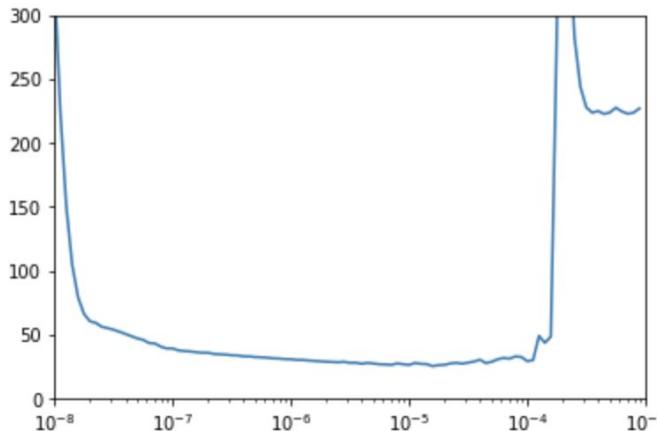


Deep Neural Networks for Time Series

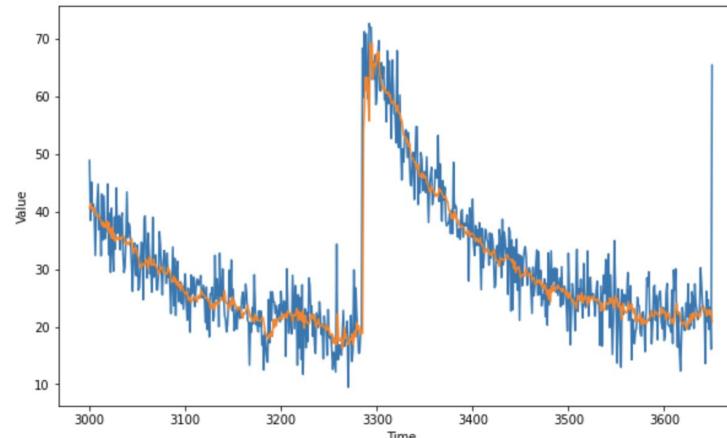
```
model = tf.keras.models.Sequential([
    tf.keras.layers.Dense(100, input_shape=[window_size], activation='relu'),
    tf.keras.layers.Dense(10, activation="relu"),
    tf.keras.layers.Dense(1)
])

model.compile(loss='mse', optimizer=tf.keras.optimizers.SGD(learning_rate=1e-6, momentum=0.9))
model.fit(dataset, epochs=my_epochs, verbose=1)
```

```
lrs = 1e-8 * (10 ** (np.arange(my_epochs) / 20))
plt.semilogx(lrs, history.history["loss"])
plt.axis([1e-8, 1e-3, 0, 300])
(1e-08, 0.001, 0.0, 300.0)
```



Total params: 3,121
Trainable params: 3,121
Non-trainable params: 0



Source: <https://www.coursera.org/learn/tensorflow-sequences-time-series-and-prediction>

Lawrence's LR optimization?

For consistent results:

```
import random  
seed = 51  
tf.random.set_seed(seed)  
random.seed = seed
```

TensorFlow > API > TensorFlow Core v2.2.0 > Python

tf.keras.callbacks.ReduceLROnPlateau

```
tf.keras.callbacks.ReduceLROnPlateau(  
    monitor='val_loss', factor=0.1, patience=10, verbose=0, mode='auto',  
    min_delta=0.0001, cooldown=0, min_lr=0, **kwargs  
)
```

Models often benefit from reducing the learning rate by a factor of 2-10 once learning stagnates. This callback monitors a quantity and if no improvement is seen for a 'patience' number of epochs, the learning rate is reduced.

Set all runs to epochs=500

Restart runtime and run all each test.

Source: https://www.tensorflow.org/api_docs/python/tf/keras/callbacks/ReduceLROnPlateau

Lawrence's LR optimization results (seed=51)

1. `optimizer=tf.keras.optimizers.SGD(lr=1e-6, momentum=0.9)`
`mae 4.500836`

2. `optimizer = tf.keras.optimizers.SGD(lr=4e-6, momentum=0.9)`
`mae 4.8579106`

3. `optimizer = tf.keras.optimizers.Adam(learning_rate=0.001)`
`mae 5.3804765`

4. `ReduceLROnPlateau(monitor='loss', patience=25, verbose=1)`
`optimizer = tf.keras.optimizers.Adam(learning_rate=0.01)`
`mae 4.490887`

Lawrence's LR optimization?

For consistent results:

```
import random  
seed = 52  
tf.random.set_seed(seed)  
random.seed = seed
```

Set all runs to epochs=500

Restart runtime and run all each test.

Lawrence's LR optimization results (seed=52)

1. `optimizer=tf.keras.optimizers.SGD(lr=1e-6, momentum=0.9)`
`mae 4.577156`

2. `optimizer = tf.keras.optimizers.SGD(lr=2e-6, momentum=0.9)`
`mae 4.5869045`

3. `optimizer = tf.keras.optimizers.Adam(learning_rate=0.001)`
`mae 4.632053`

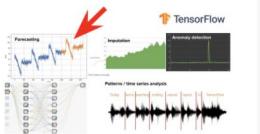
4. `ReduceLROnPlateau(monitor='loss', patience=25, verbose=1)`
`optimizer = tf.keras.optimizers.Adam(learning_rate=0.01)`
`mae 4.528657`

Let's continue our Time Series adventure 😊

FRI, JUN 19, 7:30 PM EDT

TensorFlow in Practice - C4 Week 1,2 - Forecasting and DNN for Time...

Online event



Join us for our 9th adventure in Deep Learning! Just bring your curiosity and get ready to meet our growing community 😊 We are taking Course 4 of TensorFlow in Practice Specialization available at:...



39 attendees



31 attendees

Attend

FRI, JUL 3, 7:30 PM EDT

A chat with Laurence Moroney, AI Lead at Google

Online event



Join us for a fun conversation with Laurence Moroney, AI Lead at Google (<https://www.linkedin.com/in/laurence-moroney>) and developer of our TensorFlow in Practice and TensorFlow: Data and Deployment Specializations 🎉 We plan to...



9 attendees

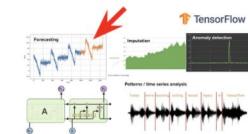
1

Attend

FRI, JUN 26, 7:30 PM EDT

TensorFlow in Practice - C4 Week 3,4 - Forecasting & RNN, Conv1D f...

Online event



Join us for our 10th adventure in Deep Learning! Just bring your curiosity and get ready to meet our growing community 😊 We are taking our last Course 4 of TensorFlow in Practice Specialization available at:...

Attend

FRI, JUL 10, 7:30 PM EDT

How to prepare for and pass the TensorFlow Developer Certificate 🎉

Online event



Join us for our 11th adventure in Deep Learning! Just bring your curiosity and get ready to meet our growing community 😊 Join Zoom Meeting: <https://us02web.zoom.us/j/84402592502?...>



14 attendees

Attend

Course 4: Sequences, Time Series and Prediction

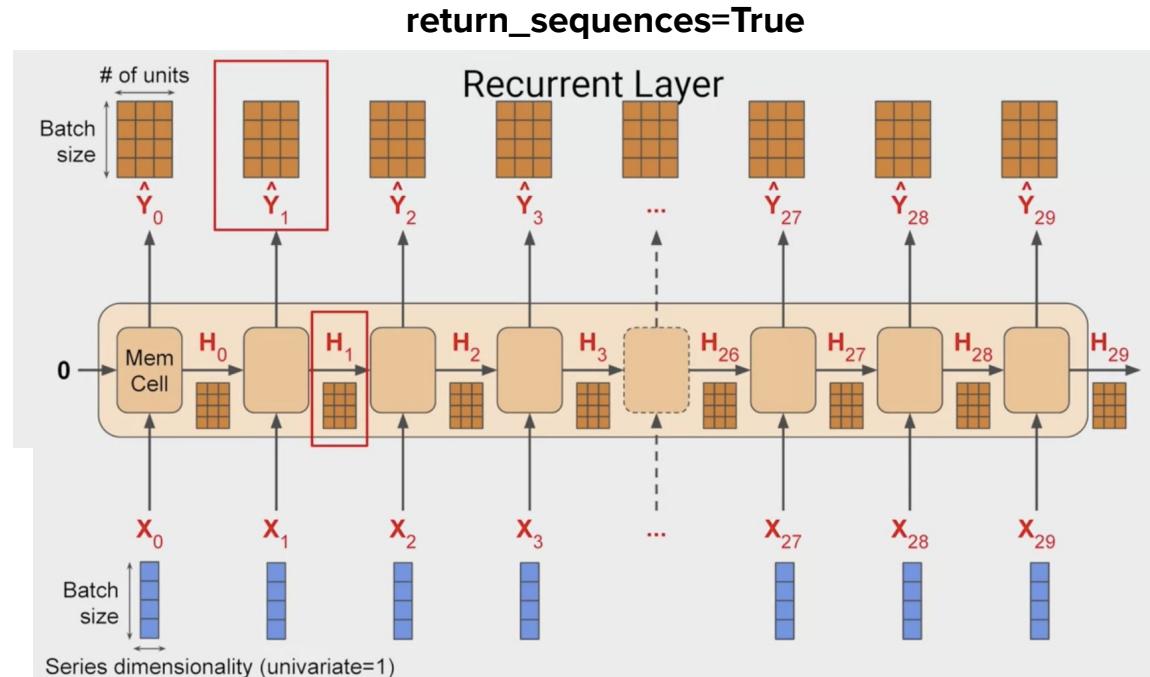
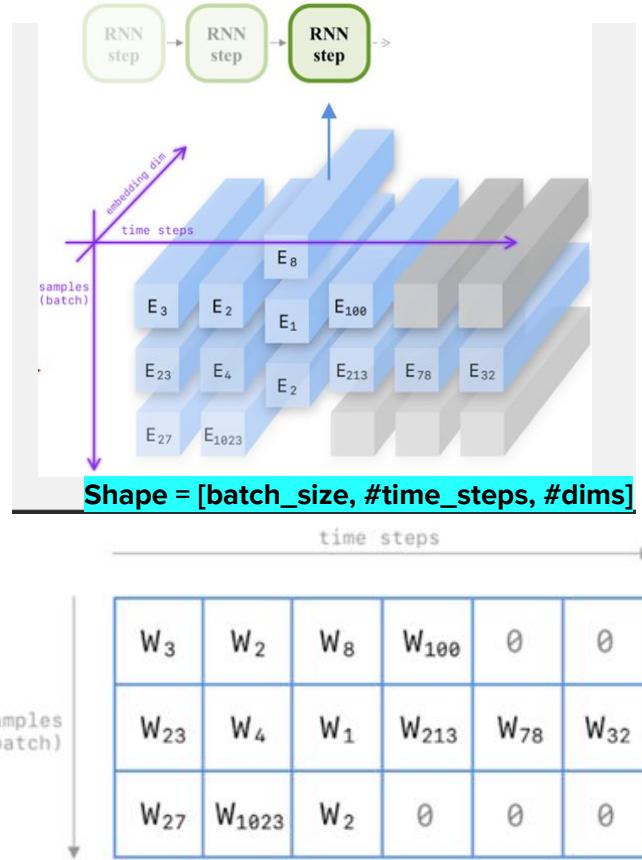
Week 1: Sequences and Prediction

Week 2: Deep Neural Networks for Time Series

Week 3: Recurrent Neural Networks for Time Series

Week 4: Real-world time series data

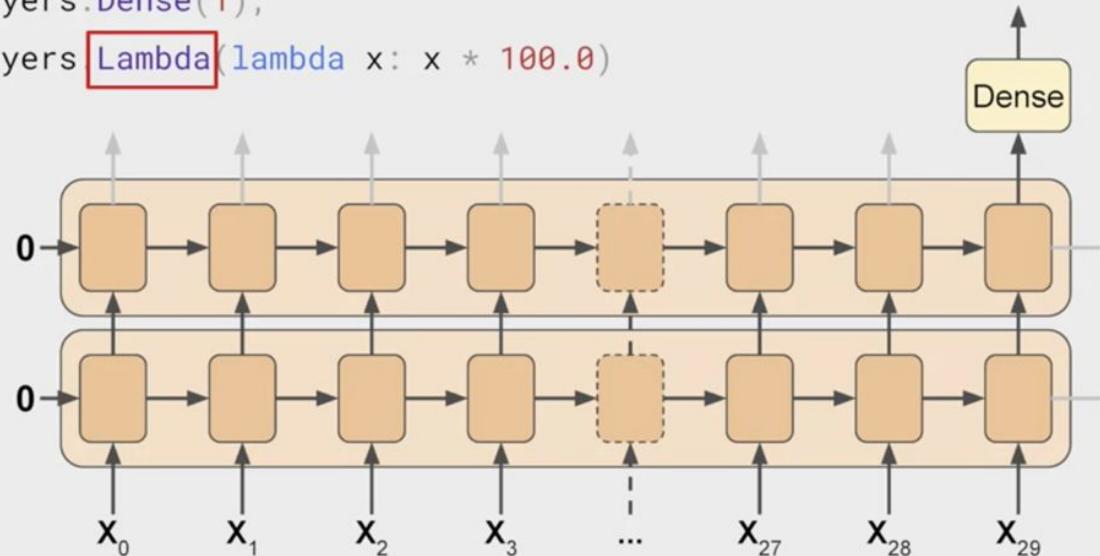
Recurrent Neural Networks for Time Series



Sources: <https://www.coursera.org/learn/tensorflow-sequences-time-series-and-prediction>
<http://karpathy.github.io/2015/05/21/rnn-effectiveness>
https://iust-deep-learning.github.io/972/static_files/assignments/assignment_05_preview.html

Recurrent Neural Networks for Time Series

```
model = keras.models.Sequential([
    keras.layers Lambda(lambda x: tf.expand_dims(x, axis=-1),
                        input_shape=[None]),
    keras.layers.SimpleRNN(20, return_sequences=True),
    keras.layers.SimpleRNN(20),
    keras.layers.Dense(1),
    keras.layers Lambda(lambda x: x * 100.0)
])
```

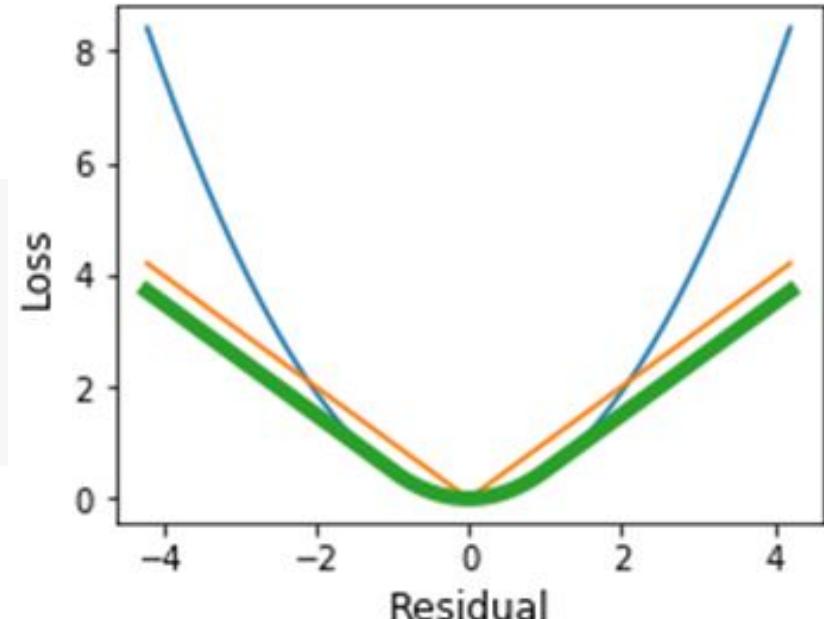


Recurrent Neural Networks for Time Series

Loss Function: Huber

$$L_\delta(y, f(x)) = \begin{cases} \frac{1}{2}(y - f(x))^2 & \text{for } |y - f(x)| \leq \delta, \\ \delta|y - f(x)| - \frac{1}{2}\delta^2 & \text{otherwise.} \end{cases}$$

```
from matplotlib import pyplot as plt
from matplotlib.pyplot import figure
import numpy as np
X=np.linspace(-4.2,4.2,841)
mse_loss = 0.5*X**2
mae_loss = np.abs(X)
huber = lambda x: (0.5*x**2 if (abs(x)<=1) else (abs(x)-0.5))
huber_loss = np.array(list(map(huber,X)))
```



```
model.compile(loss=tf.keras.losses.Huber(), ...)
```

Source: <https://www.coursera.org/learn/tensorflow-sequences-time-series-and-prediction>
https://en.wikipedia.org/wiki/Huber_loss

Recurrent Neural Networks for Time Series

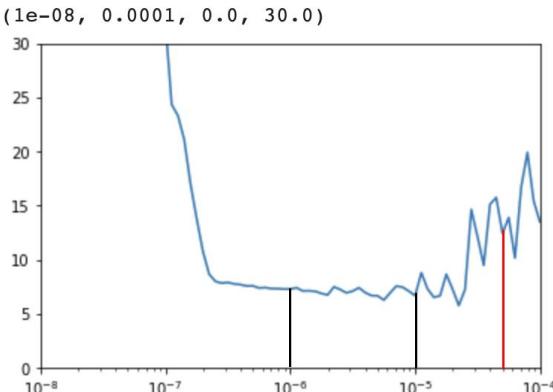
```
# Create the series
series = baseline + trend(time, slope) + seasonality(time, period=365, amplitude=10)
# Update with noise
series += noise(time, noise_level, seed=42)

split_time = 1000
time_train = time[:split_time]
x_train = series[:split_time]
time_valid = time[split_time:]
x_valid = series[split_time:]

window_size = 20
batch_size = 32
shuffle_buffer_size = 1000

my_epoch = 100

plt.semilogx(history.history["lr"], history.history["loss"])
plt.axis([1e-8, 1e-4, 0, 30])
```



```
tf.keras.backend.clear_session()
tf.random.set_seed(51)
np.random.seed(51)

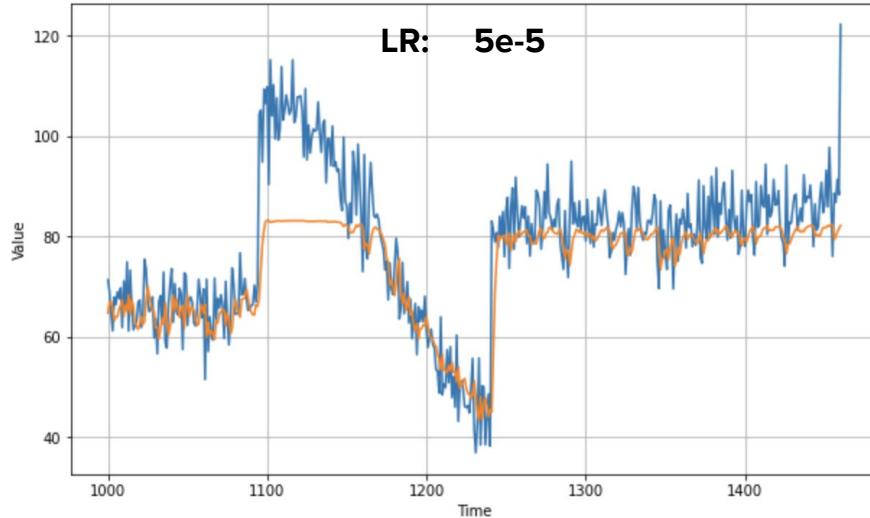
train_set = windowed_dataset(x_train, window_size, batch_size=128, shuffle_buffer=shuffle_buffer_size)

model = tf.keras.models.Sequential([
    tf.keras.layers.Lambda(lambda x: tf.expand_dims(x, axis=-1),
                           input_shape=[None]),
    tf.keras.layers.SimpleRNN(40, return_sequences=True),
    tf.keras.layers.SimpleRNN(40),
    tf.keras.layers.Dense(1),
    tf.keras.layers.Lambda(lambda x: x * 100.0)
])

lr_schedule = tf.keras.callbacks.LearningRateScheduler(
    lambda epoch: 1e-8 * 10**(epoch / 20))
optimizer = tf.keras.optimizers.SGD(lr=1e-8, momentum=0.9)
model.compile(loss=tf.keras.losses.Huber(),
              optimizer=optimizer,
              metrics=["mae"])
history = model.fit(train_set, epochs=my_epoch, callbacks=[lr_schedule])
```

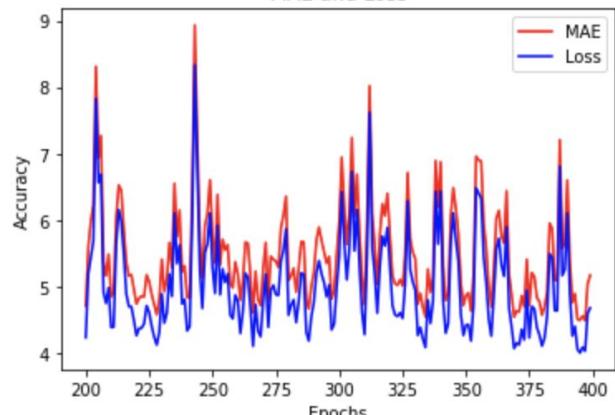
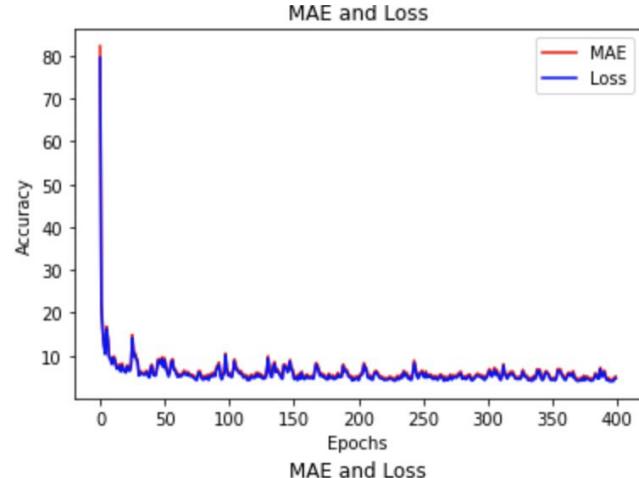
```
optimizer = tf.keras.optimizers.SGD(lr=5e-5, momentum=0.9)
model.compile(loss=tf.keras.losses.Huber(),
              optimizer=optimizer,
              metrics=["mae"])
history = model.fit(dataset, epochs=my_epoch*4)
```

Recurrent Neural Networks for Time Series

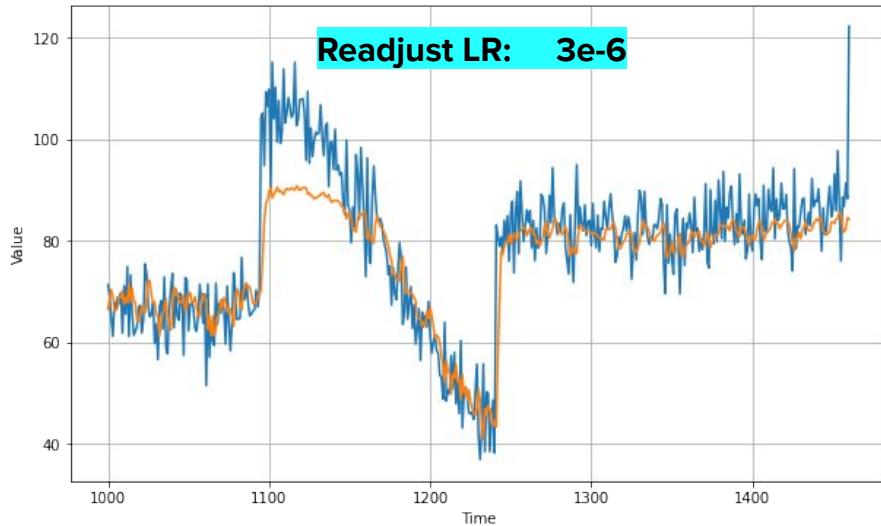


```
tf.keras.metrics.mean_absolute_error(x_valid, results).numpy()
```

6.824173

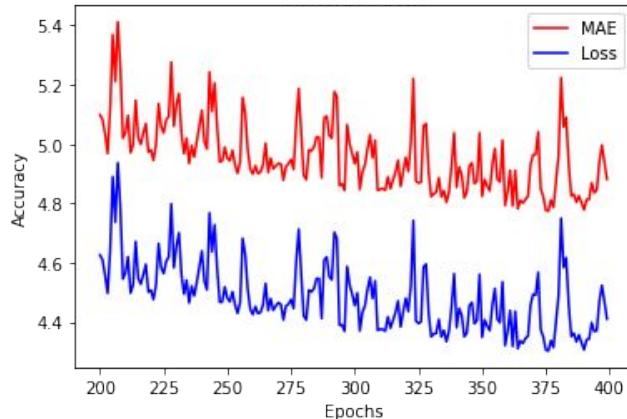
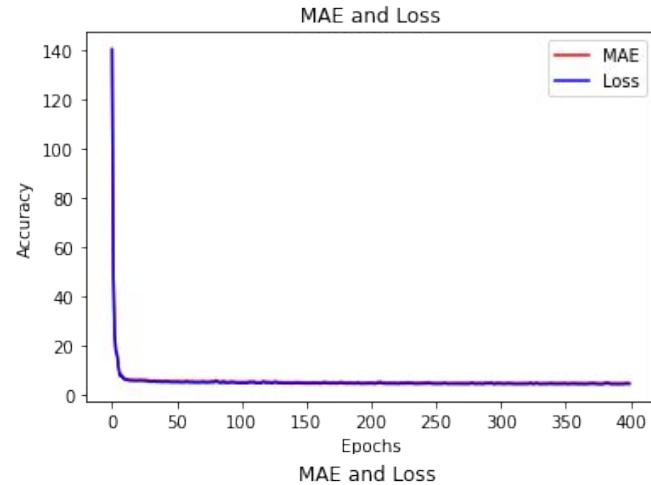


Recurrent Neural Networks for Time Series



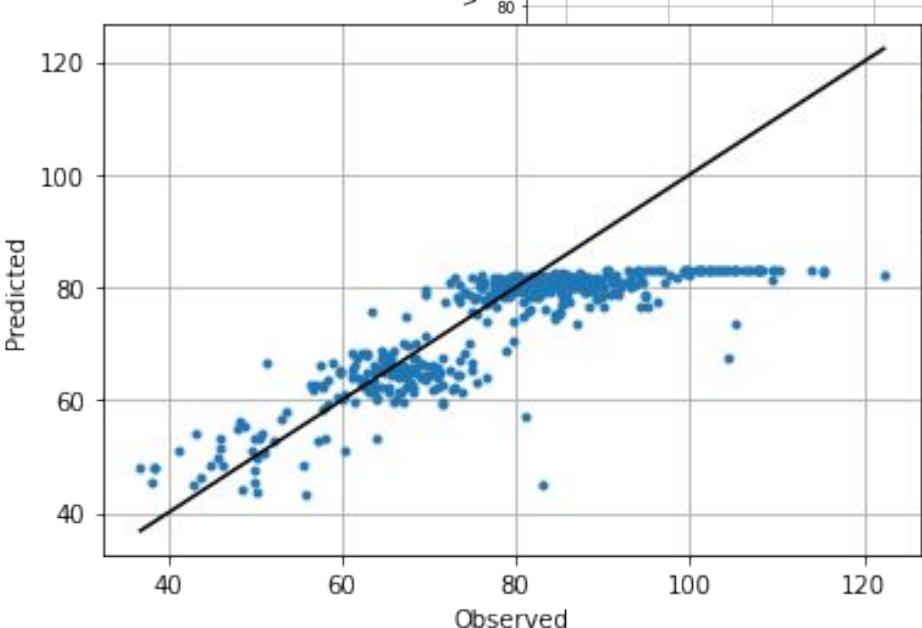
```
tf.keras.metrics.mean_absolute_error(x_valid, results).numpy()
```

5.80845

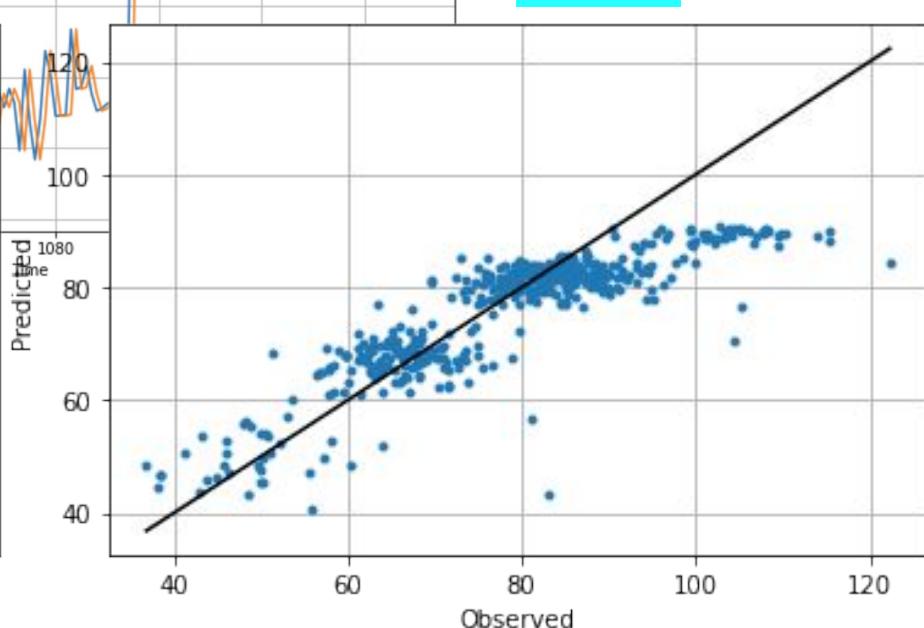


Recurrent Neural Networks for Time Series

LR: 5e-5



LR: 3e-6



Recurrent Neural Networks for Time Series

```
# Create the series
series = baseline + trend(time, slope) + seasonality(time, period=365, amplitude=amplitude)
# Update with noise
series += noise(time, noise_level, seed=42)

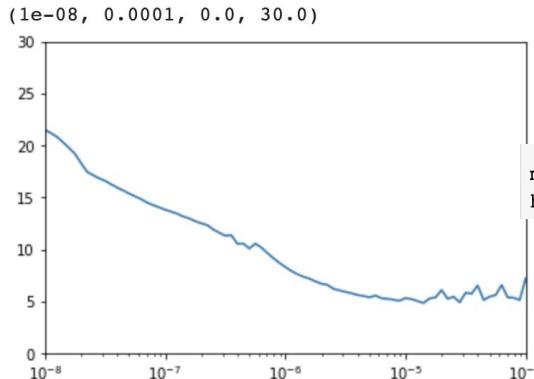
split_time = 1000
time_train = time[:split_time]
x_train = series[:split_time]
time_valid = time[split_time:]
x_valid = series[split_time:]

window_size = 20
batch_size = 32
shuffle_buffer_size = 1000

my_epoch = 100
```

time = np.arange(4 * 365 + 1,
baseline = 10
series = trend(time, 0.1)
baseline = 10
amplitude = 40
slope = 0.05
noise_level = 5

<https://> plt.semilogx(history.history["lr"], history.history["loss"])



```
tf.keras.backend.clear_session()
tf.random.set_seed(51)
np.random.seed(51)

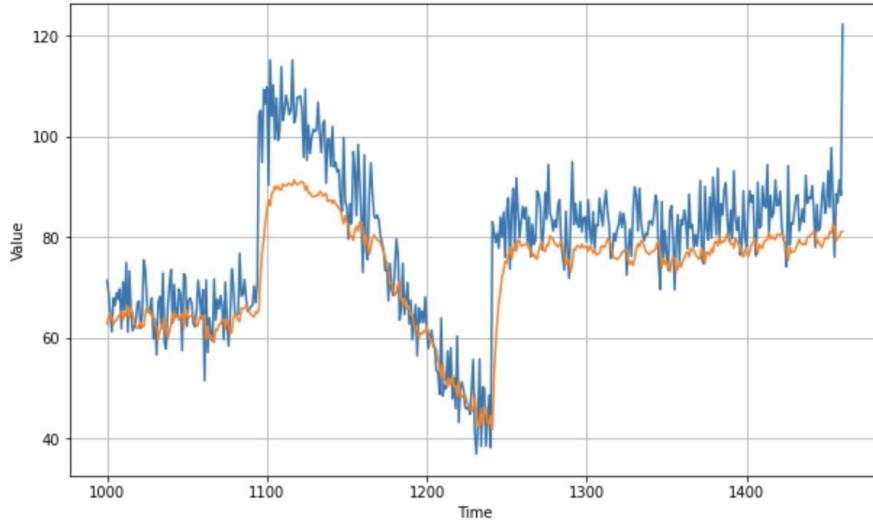
tf.keras.backend.clear_session()
dataset = windowed_dataset(x_train, window_size, batch_size, shuffle_buffer_size)

model = tf.keras.models.Sequential([
    tf.keras.layers.Lambda(lambda x: tf.expand_dims(x, axis=-1), input_shape=[None]),
    tf.keras.layers.Bidirectional(tf.keras.layers.LSTM(32, return_sequences=True)),
    tf.keras.layers.Bidirectional(tf.keras.layers.LSTM(32)),
    tf.keras.layers.Dense(1),
    tf.keras.layers.Lambda(lambda x: x * 100.0)
])

lr_schedule = tf.keras.callbacks.LearningRateScheduler(
    lambda epoch: 1e-8 * 10**((epoch / 20))
optimizer = tf.keras.optimizers.SGD(lr=1e-8, momentum=0.9)
model.compile(loss=tf.keras.losses.Huber(),
              optimizer=optimizer,
              metrics=["mae"])
history = model.fit(dataset, epochs=100, callbacks=[lr_schedule])
```

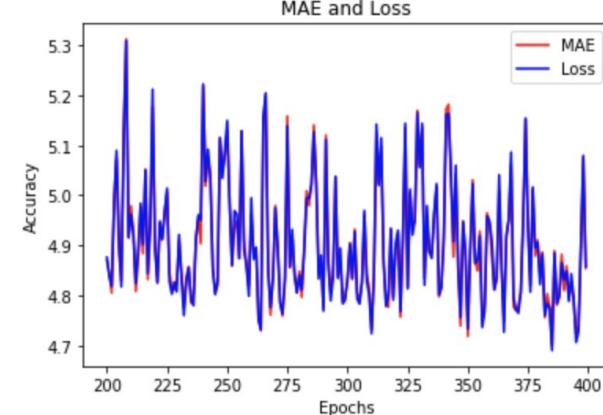
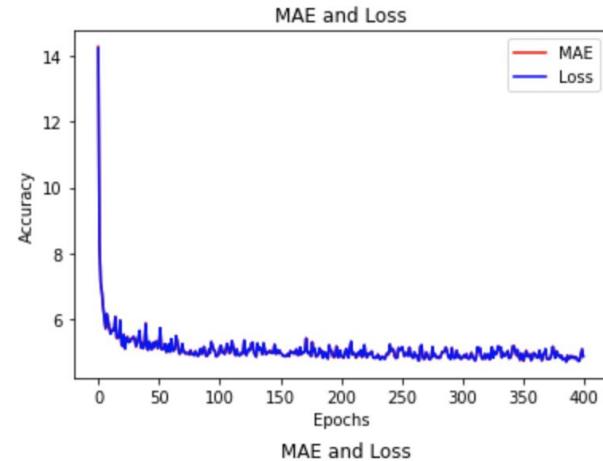
```
model.compile(loss="mae", optimizer=tf.keras.optimizers.SGD(lr=1e-5, momentum=0.9),metrics=["mae"])
history = model.fit(dataset,epochs=400,verbose=0)
```

Recurrent Neural Networks for Time Series



```
tf.keras.metrics.mean_absolute_error(x_valid, results).numpy()
```

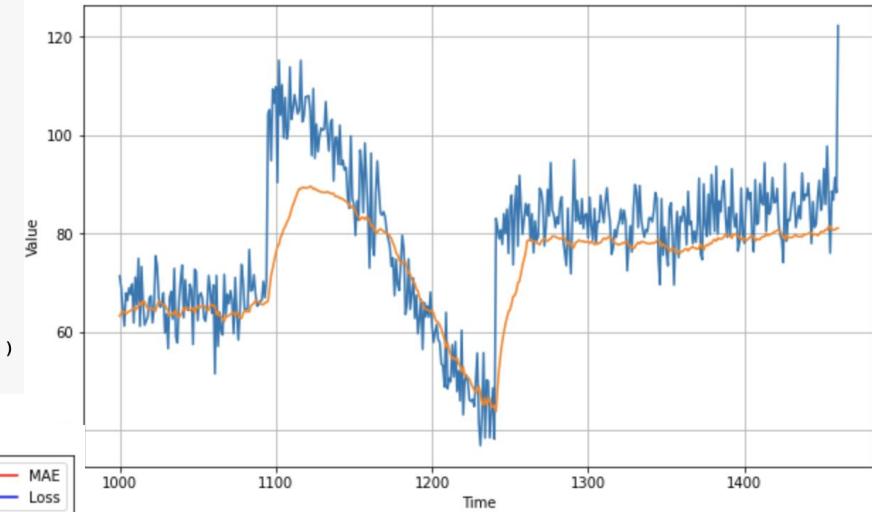
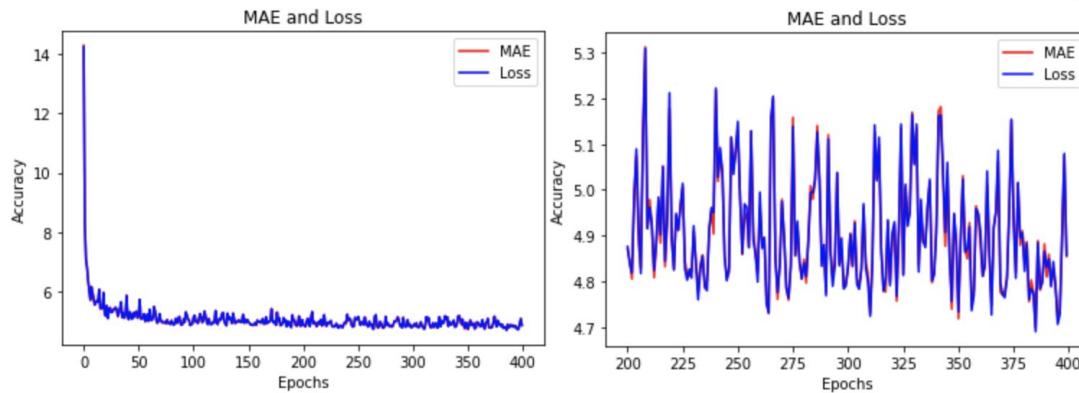
6.95316466.9531646



Source: <https://www.coursera.org/learn/tensorflow-sequences-time-series-and-prediction>

Recurrent Neural Networks for Time Series

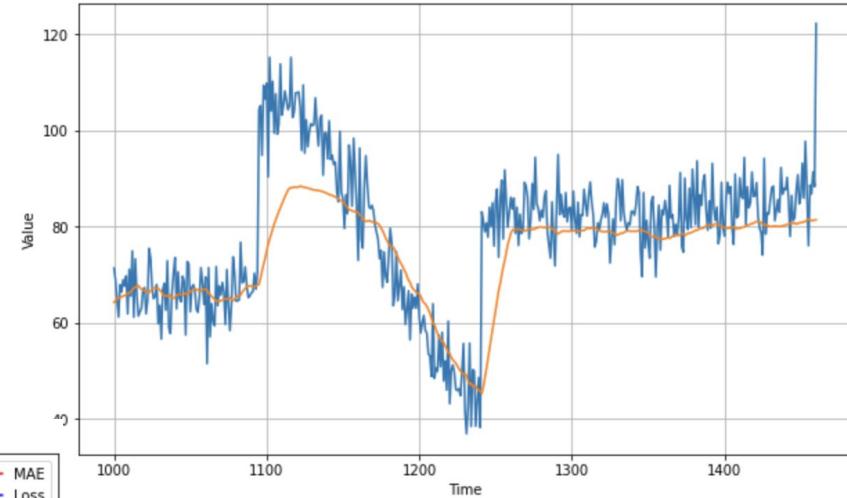
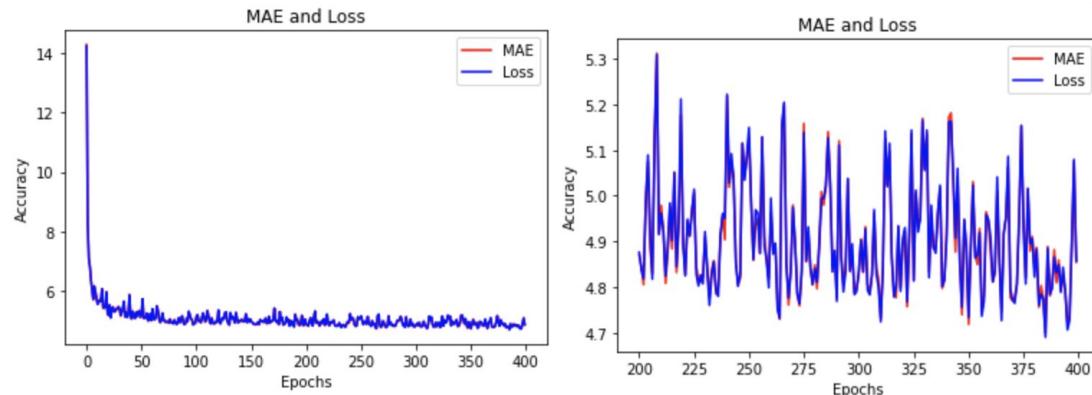
```
tf.keras.backend.clear_session()  
dataset = windowed_dataset(x_train, window_size, batch_size, shuffle_buffer_size)  
  
model = tf.keras.models.Sequential([  
    tf.keras.layers.Lambda(lambda x: tf.expand_dims(x, axis=-1),  
                           input_shape=[None]),  
    tf.keras.layers.LSTM(32, return_sequences=True),  
    tf.keras.layers.LSTM(32),  
    tf.keras.layers.Dense(1),  
    tf.keras.layers.Lambda(lambda x: x * 100.0)  
])  
  
model.compile(loss="mae", optimizer=tf.keras.optimizers.SGD(lr=1e-6, momentum=0.9))  
history = model.fit(dataset, epochs=100, verbose=0)
```



```
tf.keras.metrics.mean_absolute_error(x_valid, results).numpy()  
7.2298956
```

Recurrent Neural Networks for Time Series

```
tf.keras.backend.clear_session()  
dataset = windowed_dataset(x_train, window_size, batch_size, shuffle_buffer_size)  
  
model = tf.keras.models.Sequential([  
    tf.keras.layers.Lambda(lambda x: tf.expand_dims(x, axis=-1),  
                           input_shape=[None]),  
    tf.keras.layers.Bidirectional(tf.keras.layers.LSTM(32, return_sequences=True)),  
    tf.keras.layers.Bidirectional(tf.keras.layers.LSTM(32, return_sequences=True)),  
    tf.keras.layers.Bidirectional(tf.keras.layers.LSTM(32)),  
    tf.keras.layers.Dense(1),  
    tf.keras.layers.Lambda(lambda x: x * 100.0)  
])  
  
model.compile(loss="mae", optimizer=tf.keras.optimizers.SGD(lr=1e-6, momentum=0.9))  
model.fit(dataset, epochs=100)
```

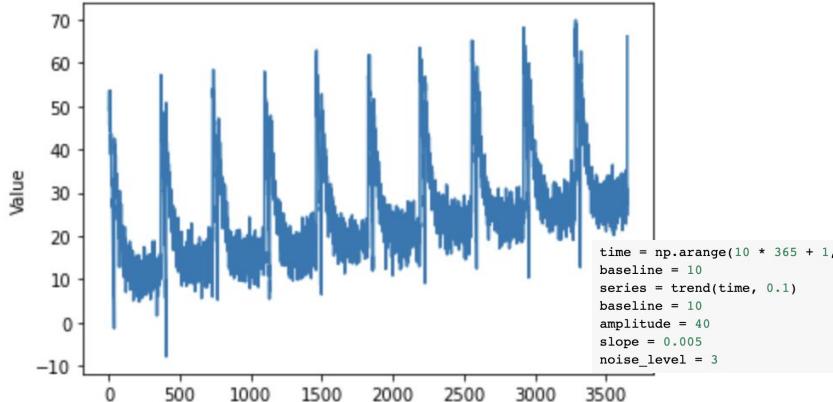


```
tf.keras.metrics.mean_absolute_error(x_valid, results).numpy()
```

7.212102

```
epochs_comment = """  
Lambda, 2X Bi LSTM(32), 400 epochs  
6.7509785  
Lambda, 2X LSTM(32), 100 epochs  
7.2298923  
Lambda, 3X Bi LSTM(32), 100 epochs  
7.242644  
"""
```

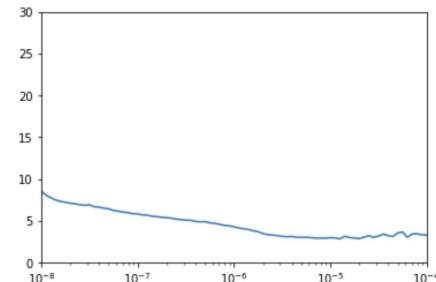
Recurrent Neural Networks for Time Series



```
plt.semilogx(history.history["lr"], history.history["loss"])
plt.axis([1e-08, 1e-04, 0, 30])
```

FROM THIS PICK A LEARNING RATE

```
(1e-08, 0.0001, 0.0, 30.0)
```



```
tf.keras.backend.clear_session()
tf.random.set_seed(51)
np.random.seed(51)

tf.keras.backend.clear_session()
dataset = windowed_dataset(x_train, window_size, batch_size, shuffle_buffer_size)

model = tf.keras.models.Sequential([
    tf.keras.layers.Lambda(lambda x: tf.expand_dims(x, axis=-1), input_shape=[None]),
    tf.keras.layers.Bidirectional(tf.keras.layers.LSTM(32, return_sequences=True)),
    tf.keras.layers.Bidirectional(tf.keras.layers.LSTM(32)),
    tf.keras.layers.Dense(1),
    tf.keras.layers.Lambda(lambda x: x * 100.0)
])

lr_schedule = tf.keras.callbacks.LearningRateScheduler(lambda epoch: 1e-8 * 10**((epoch / 20)))

optimizer = tf.keras.optimizers.SGD(lr=1e-8, momentum=0.9)

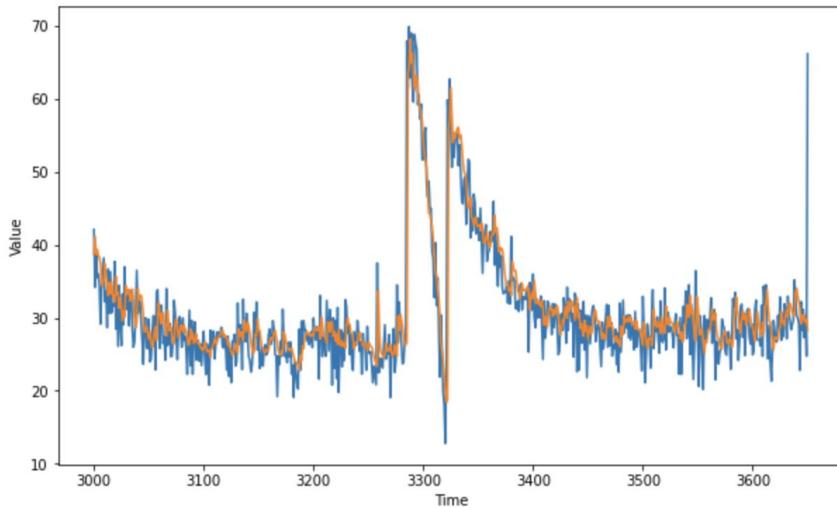
model.compile(loss=tf.keras.losses.Huber(),
              optimizer=optimizer,
              metrics=["mae"])

history = model.fit(dataset, epochs=100, callbacks=[lr_schedule])
```



```
model.compile(loss="mse", optimizer=tf.keras.optimizers.SGD(lr=1e-5, momentum=0.9),metrics=["mae"])
history = model.fit(dataset,epochs=400,verbose=1)
```

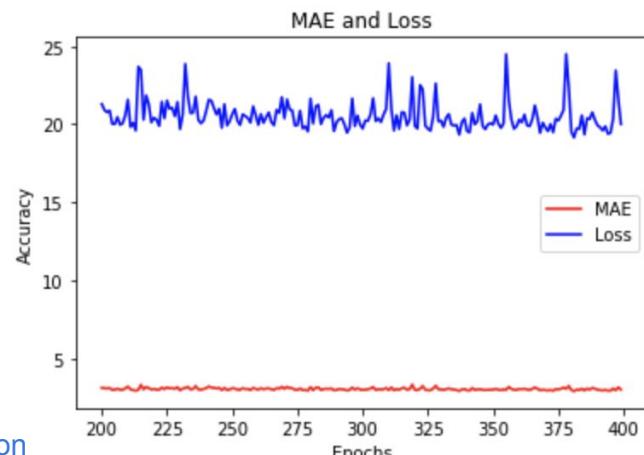
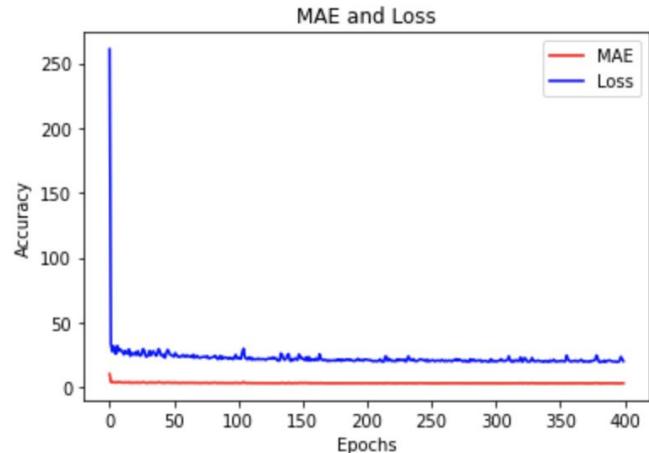
Recurrent Neural Networks for Time Series



```
tf.keras.metrics.mean_absolute_error(x_valid, results).numpy()
```

```
# YOUR RESULT HERE SHOULD BE LESS THAN 4
```

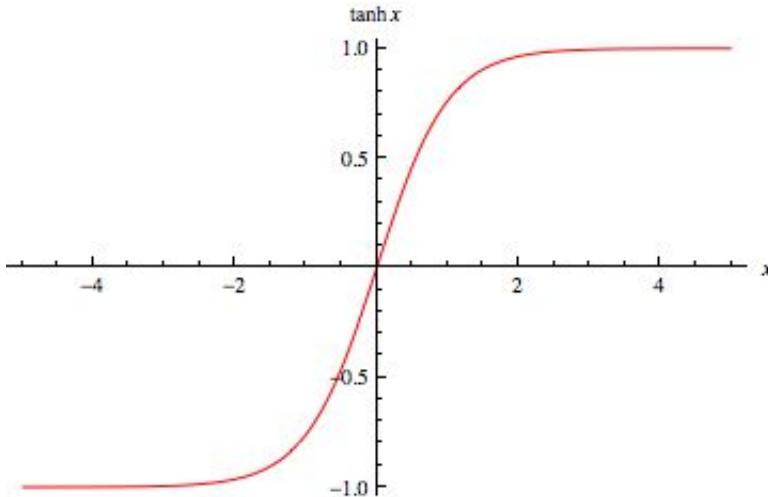
```
3.051421
```



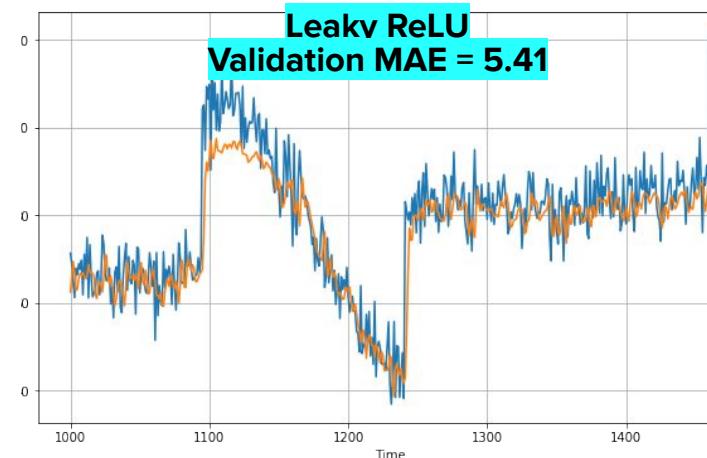
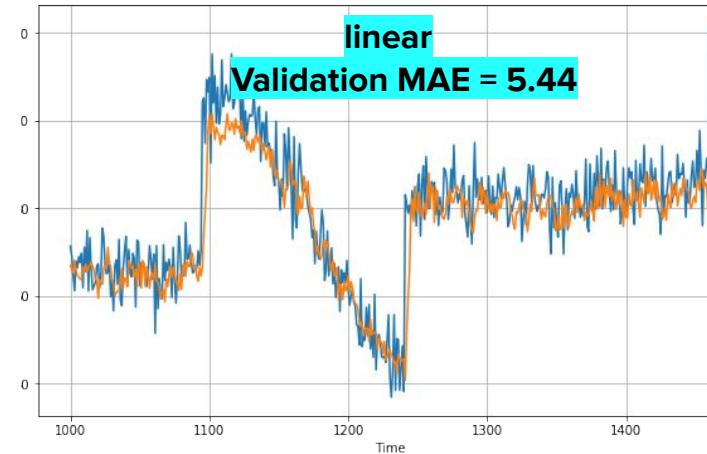
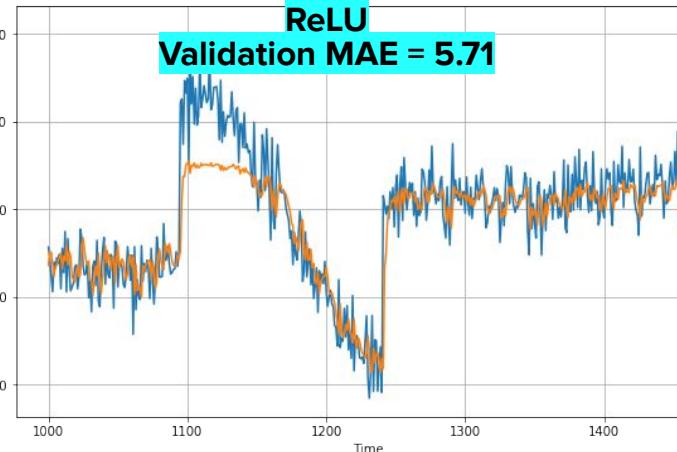
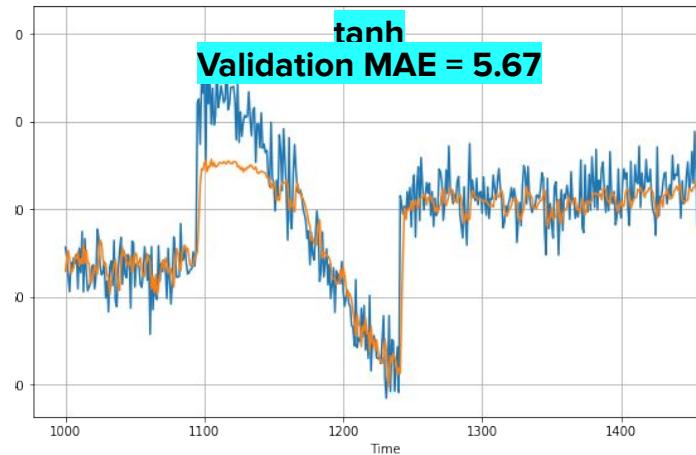
Recurrent Neural Networks for Time Series



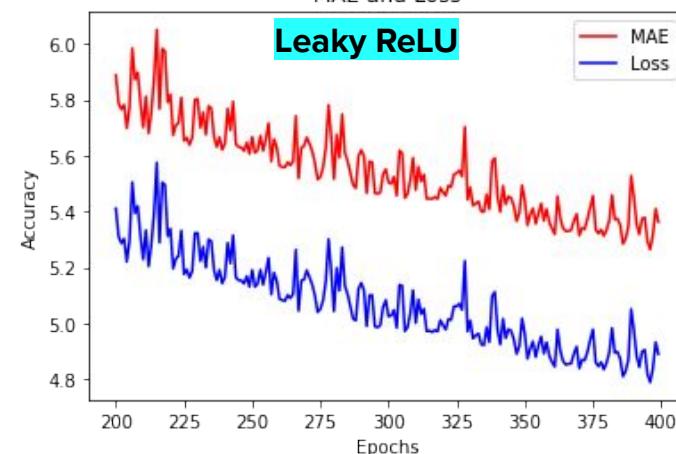
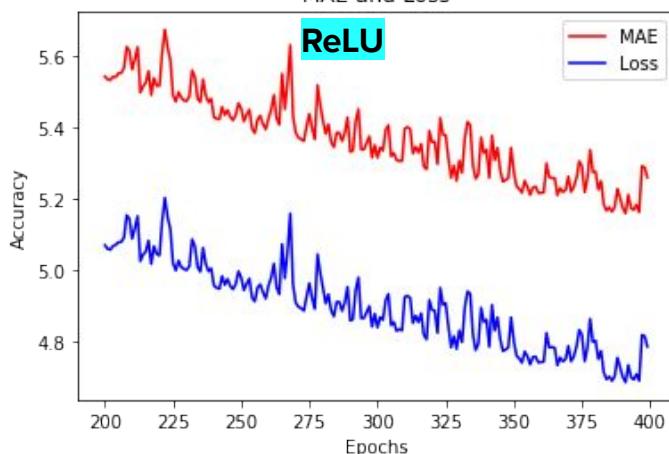
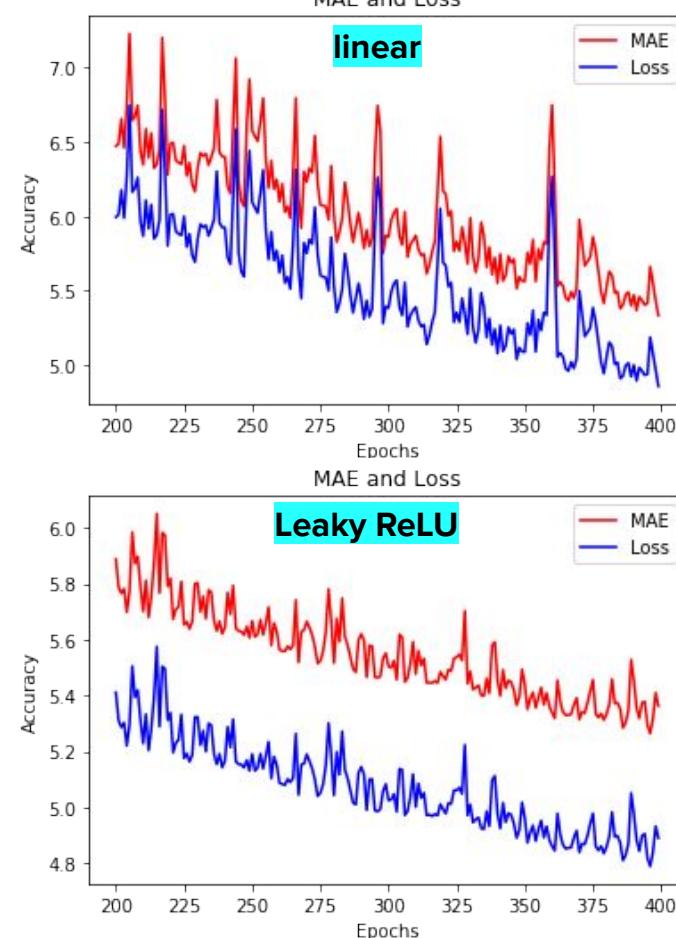
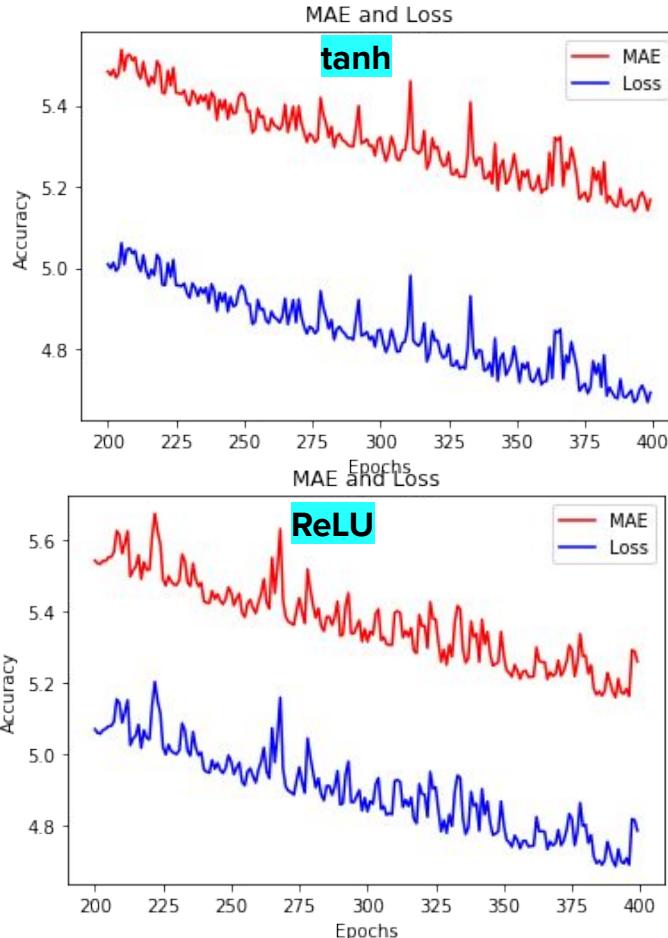
```
tf.keras.layers.SimpleRNN(  
    units, activation='tanh', use_bias=True, kernel_initializer='glor  
    recurrent_initializer='orthogonal', bias_initializer='zeros',  
    kernel_regularizer=None, recurrent_regularizer=None, bias_regularizer  
    activity_regularizer=None, kernel_constraint=None, recurrent_cons
```



Recurrent Neural Networks for Time Series



Recurrent Neural Networks for Time Series

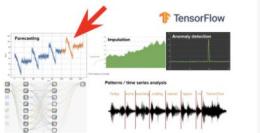


Let's continue our Time Series adventure 😊

FRI, JUN 19, 7:30 PM EDT

TensorFlow in Practice - C4 Week 1,2 - Forecasting and DNN for Time...

Online event



Join us for our 9th adventure in Deep Learning! Just bring your curiosity and get ready to meet our growing community 😊 We are taking Course 4 of TensorFlow in Practice Specialization available at:...



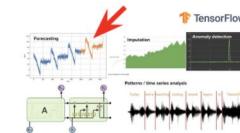
39 attendees



FRI, JUN 26, 7:30 PM EDT

TensorFlow in Practice - C4 Week 3,4 - Forecasting & RNN, Conv1D f...

Online event



Join us for our 10th adventure in Deep Learning! Just bring your curiosity and get ready to meet our growing community 😊 We are taking our last Course 4 of TensorFlow in Practice Specialization available at:...

[Attend](#)

FRI, JUL 3, 7:30 PM EDT

A chat with Laurence Moroney, AI Lead at Google

Online event



Join us for a fun conversation with Laurence Moroney, AI Lead at Google (<https://www.linkedin.com/in/laurence-moroney>) and developer of our TensorFlow in Practice and TensorFlow: Data and Deployment Specializations 🎉 We plan to...



9 attendees



1

[Attend](#)

FRI, JUL 10, 7:30 PM EDT

How to prepare for and pass the TensorFlow Developer Certificate 🎉

Online event



Join us for our 11th adventure in Deep Learning! Just bring your curiosity and get ready to meet our growing community 😊 Join Zoom Meeting: <https://us02web.zoom.us/j/84402592502?...>



14 attendees

[Attend](#)

Course 4: Sequences, Time Series and Prediction

Week 1: Sequences and Prediction

Week 2: Deep Neural Networks for Time Series

Week 3: Recurrent Neural Networks for Time Series

Week 4: Real-world time series data

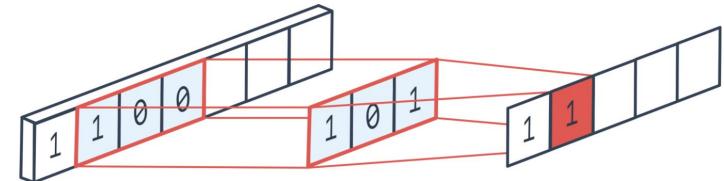
Real-world time series data

TensorFlow > API > TensorFlow Core v2.2.0 > Python

tf.keras.layers.Conv1D

1D convolution layer (e.g. temporal convolution).

```
tf.keras.layers.Conv1D(  
    filters, kernel_size, strides=1, padding='valid', data_format='channels_last',  
    dilation_rate=1, activation=None, use_bias=True,  
    kernel_initializer='glorot_uniform', bias_initializer='zeros',  
    kernel_regularizer=None, bias_regularizer=None, activity_regularizer=None,  
    kernel_constraint=None, bias_constraint=None, **kwargs  
)  
  
>>> # The inputs are 128-length vectors with 10 timesteps, and the batch size  
>>> # is 4.  
>>> input_shape = (4, 10, 128)  
>>> x = tf.random.normal(input_shape)  
>>> y = tf.keras.layers.Conv1D(  
... 32, 3, activation='relu', input_shape=input_shape)(x)  
>>> print(y.shape)  
(4, 8, 32)
```



Conv1d-Input1d Example [Image [12]]

filters

kernel_size

strides

Integer, the dimensionality of the output space (i.e. the number of output filters in the convolution).

An integer or tuple/list of a single integer, specifying the length of the 1D convolution window.

An integer or tuple/list of a single integer, specifying the stride length of the convolution. Specifying any stride value != 1 is incompatible with specifying any `dilation_rate` value != 1.

Input shape:

3D tensor with shape: (batch_size, steps, input_dim)

Output shape:

3D tensor with shape: (batch_size, new_steps, filters) `steps` value might have changed due to padding or strides.

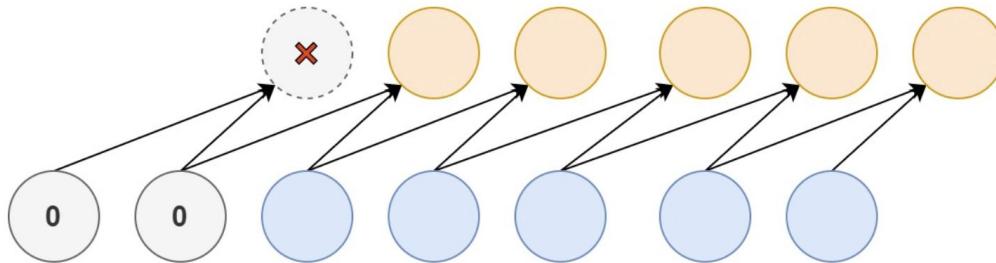
Source: https://www.tensorflow.org/api_docs/python/tf/keras/layers/Conv1D

<https://towardsdatascience.com/pytorch-basics-how-to-train-your-neural-net-intro-to-cnn-26a14c2ea29>

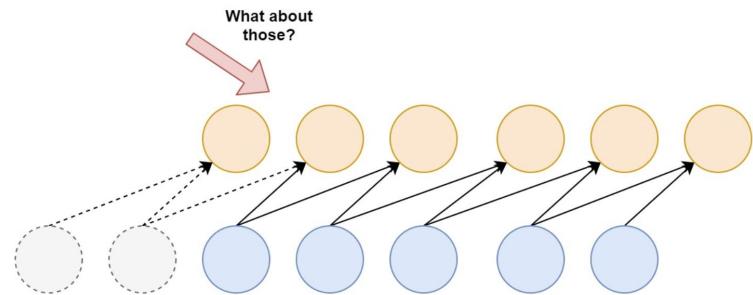
Real-world time series data

layer using "causal" padding: this ensures that the convolutional layer does not peek into the future when making predictions (it is equivalent to padding the inputs with the right amount of zeros on the left and using "valid" padding). We then add similar

Causal padding on the Conv1D layer allows you to include the partial information in your training process. By padding your input dataset with zeros at the front, a causal mapping to the first, missed-out targets can be made (Keras, n.d.; The Blog, n.d.). While the first target will be useless for training, the second can now be used based on the partial information that we have:



But what about the first two targets?



Although they are valid targets, the *inputs* are incomplete – that is, there is insufficient input data available in order to successfully use them in the training process (The Blog, n.d.). For the second target, *one* input – visible in gray – is missing (whereas the second is actually there), while for the first target both aren't there.

For the first target, there is no real hope for success (as we don't have any input at all and hence do not know which values produce the target value), but for the second, we have a partial picture: we've got half the inputs that produce the target.

Source: <https://books.google.com/books?id=HnetDwAAQBAJ>

<https://www.machinecurve.com/index.php/2020/02/07/what-is-padding-in-a-neural-network/#causal-padding>

Real-world time series d

Dilated and causal convolution

As discussed in the section on backtesting, we have to make sure that our model does not have bias:

As the convolution ensures that it

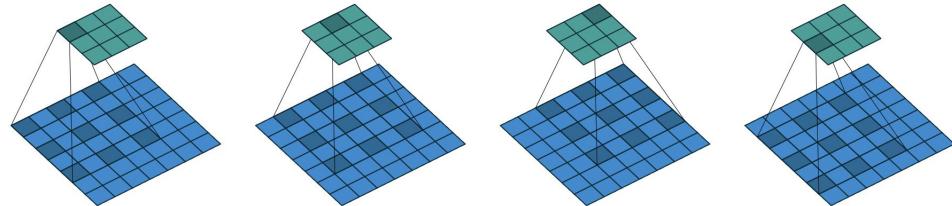
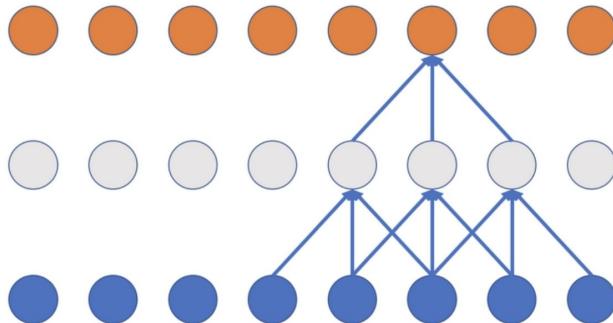


Figure 5.1: (Dilated convolution) Convolving a 3×3 kernel over a 7×7 input with a dilation factor of 2 (i.e., $i = 7$, $k = 3$, $d = 2$, $s = 1$ and $p = 0$).

Standard Convolution

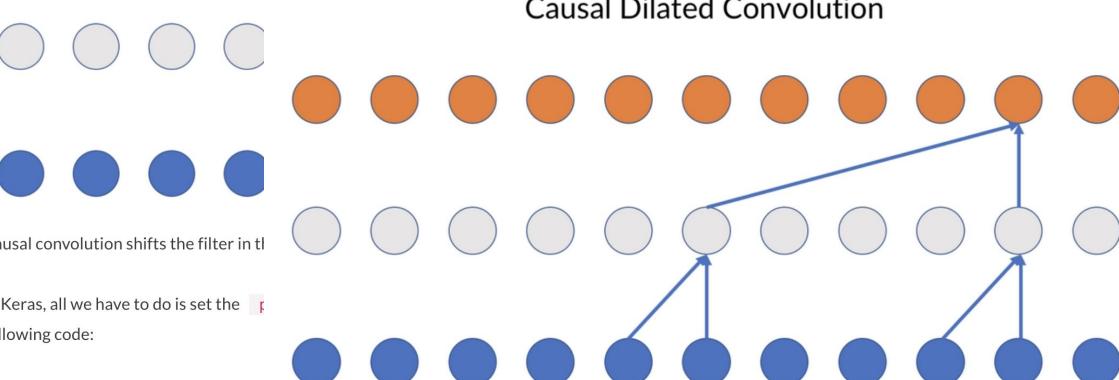


Standard convolution does not take the direction of convolution into account

As the convolutional filter slides over the data, it looks into the future as well as the past.

Dilated convolution skips over inputs while convolving

Causal Dilated Convolution



Causal convolution shifts the filter in time

In Keras, all we have to do is set the padding='causal' parameter in the following code:

Source: https://subscription.packtpub.com/book/machine_learning/9781789136364/4/ch04lvl1sec59/dilated-and-causal-convolution
<https://theblog.github.io/post/convolution-in-autoregressive-neural-networks>
<https://arxiv.org/pdf/1603.07285.pdf>

Real-world time series data

```
time = np.arange(4 * 365 + 1, dtype="float32")
baseline = 10
series = trend(time, 0.1)
baseline = 10
amplitude = 40
slope = 0.05
noise_level = 5

# Create the series
series = baseline + trend(time, slope) + seasonality(time, period=365, amplitude=amplitude)
# Update with noise
series += noise(time, noise_level, seed=42)

split_time = 1000
time_train = time[:split_time]
x_train = series[:split_time]
time_valid = time[split_time:]
x_valid = series[split_time:]

window_size = 20
batch_size = 32
shuffle_buffer_size = 1000
```

```
model.summary()

Model: "sequential"

Layer (type)          Output Shape       Param #
=====
conv1d (Conv1D)        (None, None, 32)    192
bidirectional (Bidirectional) (None, None, 64) 16640
bidirectional_1 (Bidirection) (None, None, 64) 24832
dense (Dense)          (None, None, 1)      65
lambda (Lambda)         (None, None, 1)      0
=====
Total params: 41,729
Trainable params: 41,729
Non-trainable params: 0
```

```
history = model.fit(train_set, epochs=100, callbacks=[lr_schedule])
```

```
tf.keras.backend.clear_session()
tf.random.set_seed(51)
np.random.seed(51)

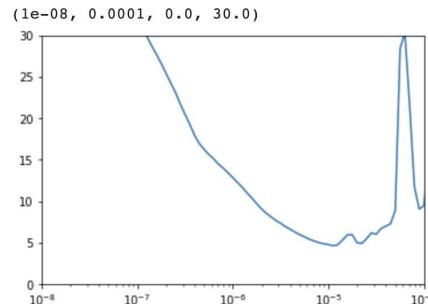
window_size = 30
train_set = windowed_dataset(x_train, window_size, batch_size=128, shuffle_buffer=shuffle_buffer_size)

model = tf.keras.models.Sequential([
    tf.keras.layers.Conv1D(filters=32, kernel_size=5,
                          strides=1, padding="causal",
                          activation="relu",
                          input_shape=[None, 1]),
    tf.keras.layers.Bidirectional(tf.keras.layers.LSTM(32, return_sequences=True)),
    tf.keras.layers.Bidirectional(tf.keras.layers.LSTM(32, return_sequences=True)),
    tf.keras.layers.Dense(1),
    tf.keras.layers.Lambda(lambda x: x * 200)
])
lr_schedule = tf.keras.callbacks.LearningRateScheduler(lambda epoch: 1e-8 * 10**(epoch / 20))

optimizer = tf.keras.optimizers.SGD(lr=1e-8, momentum=0.9)

model.compile(loss=tf.keras.losses.Huber(),
              optimizer=optimizer,
              metrics=["mae"])

plt.semilogx(history.history["lr"], history.history["loss"])
plt.axis([1e-8, 1e-4, 0, 30])
```



Source: <https://www.coursera.org/learn/tensorflow-sequences-time-series-and-prediction>

Real-world time series data

```
tf.keras.backend.clear_session()
tf.random.set_seed(51)
np.random.seed(51)
#batch_size = 16
dataset = windowed_dataset(x_train, window_size, batch_size, shuffle_buffer_size)

model = tf.keras.models.Sequential([
    tf.keras.layers.Conv1D(filters=32, kernel_size=3,
                          strides=1, padding="causal",
                          activation="relu",
                          input_shape=[None, 1]),
    tf.keras.layers.LSTM(32, return_sequences=True),
    tf.keras.layers.LSTM(32, return_sequences=True),
    tf.keras.layers.Dense(1),
    tf.keras.layers.Lambda(lambda x: x * 200)
])

optimizer = tf.keras.optimizers.SGD(lr=1e-5, momentum=0.9)

model.compile(loss=tf.keras.losses.Huber(),
              optimizer=optimizer,
              metrics=[ "mae"])

rnn_forecast = model_forecast(model, series[..., np.newaxis], window_size)

rnn_forecast.shape

(1432, 30, 1)
```

```
model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv1d (Conv1D)	(None, None, 32)	128
lstm (LSTM)	(None, None, 32)	8320
lstm_1 (LSTM)	(None, None, 32)	8320
dense (Dense)	(None, None, 1)	33
lambda (Lambda)	(None, None, 1)	0
Total params: 16,801		
Trainable params: 16,801		
Non-trainable params: 0		

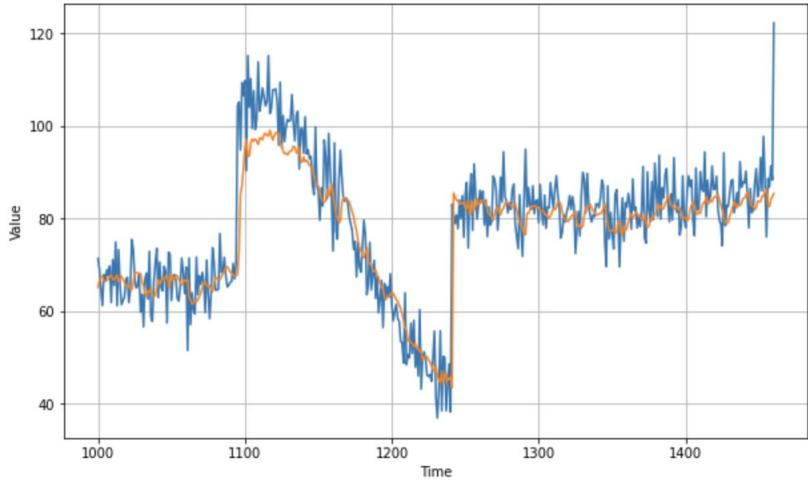
```
history = model.fit(dataset, epochs=500)
```

```
rnn_forecast = rnn_forecast[split_time - window_size:-1, -1, 0]
```

```
rnn_forecast.shape
```

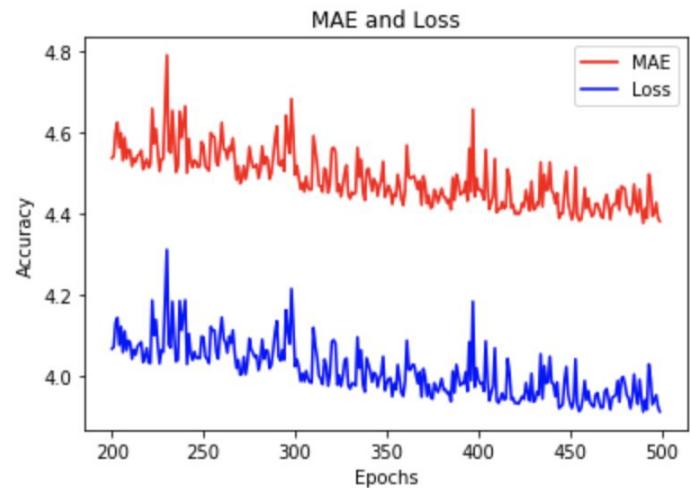
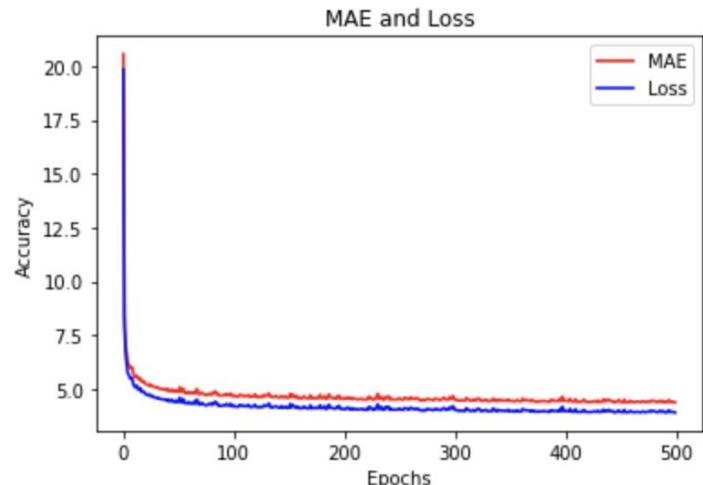
```
(461,)
```

Real-world time series data



```
tf.keras.metrics.mean_absolute_error(x_valid, rnn_forecast).numpy()
```

5.091009



Real-world time series data

Dataset

Sunspots

Monthly Mean Total Sunspot Number - from 1749 to July 2018

Especuloide • updated 5 months ago (Version 2)

Data Tasks Kernels (19) Discussion Activity Metadata

Download (70 KB) New Notebook

Usability 8.2 License CC0: Public Domain Tags natural and physical sciences, astronomy

Description

Context

Sunspots are temporary phenomena on the Sun's photosphere that appear as spots darker than the surrounding areas. They are regions of reduced surface temperature caused by concentrations of magnetic field flux that inhibit convection. Sunspots usually appear in pairs of opposite magnetic polarity. Their number varies according to the approximately 11-year solar cycle.

Source: <https://en.wikipedia.org/wiki/Sunspot>

Content :

Monthly Mean Total Sunspot Number, from 1749/01/01 to 2017/08/31

< Sunspots.csv (69.52 KB)

Detail Compact Column

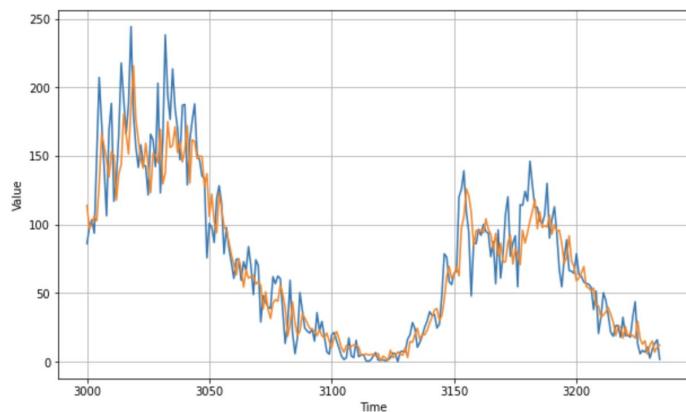
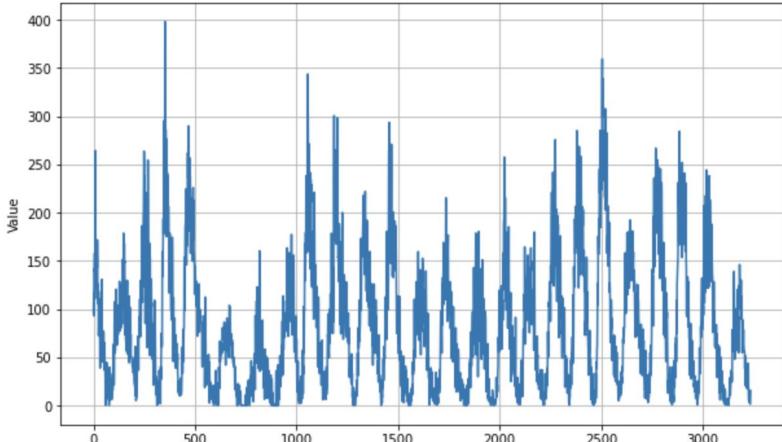
About this file

Sunspots - Monthly Mean Total Sunspot Number

#	Date	Monthly Mean T...
0	3251	398
0	1749-01-31	96.7
1	1749-02-28	104.3
2	1749-03-31	116.7
3	1749-04-30	92.8
4	1749-05-31	141.7
5	1749-06-30	139.2
6	1749-07-31	158.0
7	1749-08-31	110.5
8	1749-09-30	126.5
9	1749-10-31	125.8
10	1749-11-30	264.3

Source: <https://www.kaggle.com/robervalt/sunspots>

Real-world time series data



Source: <https://www.coursera.org/learn/tensorflow-sequences-time-series-and-prediction>

```
dataset = windowed_dataset_dnn(x_train, window_size, batch_size, shuffle_buffer_size)

model = tf.keras.models.Sequential([
    tf.keras.layers.Dense(20, input_shape=[window_size], activation="relu"),
    tf.keras.layers.Dense(10, activation="relu"),
    tf.keras.layers.Dense(1)
])

model.compile(loss="mse", optimizer=tf.keras.optimizers.SGD(lr=1e-7, momentum=0.9))

model.summary()

Model: "sequential_1"
-----  

Layer (type)           Output Shape        Param #
-----  

dense_3 (Dense)        (None, 20)          620  

dense_4 (Dense)        (None, 10)          210  

dense_5 (Dense)        (None, 1)           11  

-----  

Total params: 841  

Trainable params: 841  

Non-trainable params: 0  

-----  

model.fit(dataset, epochs=100, verbose=1)  

-----  

tf.keras.metrics.mean_absolute_error(x_valid, results).numpy()  

-----  

14.931627
```

Real-world time series data

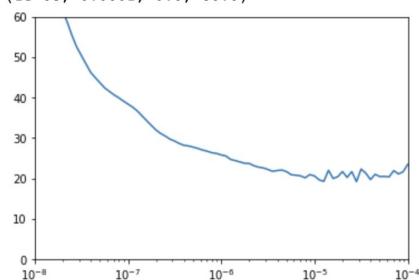
```
tf.keras.backend.clear_session()
tf.random.set_seed(51)
np.random.seed(51)
window_size = 64
batch_size = 256
train_set = windowed_dataset(x_train, window_size, batch_size, shuffle_buffer_size)
print(train_set)
print(x_train.shape)

model = tf.keras.models.Sequential([
    tf.keras.layers.Conv1D(filters=32, kernel_size=5,
                          strides=1, padding="causal",
                          activation="relu",
                          input_shape=[None, 1]),
    tf.keras.layers.LSTM(64, return_sequences=True),
    tf.keras.layers.LSTM(64, return_sequences=True),
    tf.keras.layers.Dense(30, activation="relu"),
    tf.keras.layers.Dense(10, activation="relu"),
    tf.keras.layers.Dense(1),
    tf.keras.layers.Lambda(lambda x: x * 400)
])

lr_schedule = tf.keras.callbacks.LearningRateScheduler(lambda epoch: 1e-8 * 10**((epoch / 20)))

optimizer = tf.keras.optimizers.SGD(lr=1e-8, momentum=0.9)
model.compile(loss=tf.keras.losses.Huber(),
              optimizer=optimizer,
              metrics=["mae"])

plt.semilogx(history.history["lr"], history.history["loss"])
plt.axis([1e-8, 1e-4, 0, 60])
plt.title("(1e-08, 0.0001, 0.0, 60.0)
```



```
model.summary()

<PrefetchDataset shapes: ((None, None, 1), (None, None, 1)), types:
(3000,)
Model: "sequential"

Layer (type)                 Output Shape            Param #
=====
conv1d (Conv1D)               (None, None, 32)        192
lstm (LSTM)                  (None, None, 64)       24832
lstm_1 (LSTM)                (None, None, 64)       33024
dense (Dense)                (None, None, 30)       1950
dense_1 (Dense)              (None, None, 10)      310
dense_2 (Dense)              (None, None, 1)        11
lambda (Lambda)              (None, None, 1)        0
=====
Total params: 60,319
Trainable params: 60,319
Non-trainable params: 0
```

```
history = model.fit(train_set, epochs=100, callbacks=[lr_schedule])
```

Real-world time series data

```
tf.keras.backend.clear_session()
tf.random.set_seed(51)
np.random.seed(51)
train_set = windowed_dataset(x_train, window_size=60, batch_size=100, shuffle_buffer=shuffle_buffer_size)
model = tf.keras.models.Sequential([
    tf.keras.layers.Conv1D(filters=60, kernel_size=5,
                          strides=1, padding="causal",
                          activation="relu",
                          input_shape=[None, 1]),
    tf.keras.layers.LSTM(60, return_sequences=True),
    tf.keras.layers.LSTM(60, return_sequences=True),
    tf.keras.layers.Dense(30, activation="relu"),
    tf.keras.layers.Dense(10, activation="relu"),
    tf.keras.layers.Dense(1),
    tf.keras.layers.Lambda(lambda x: x * 400)
])

optimizer = tf.keras.optimizers.SGD(lr=1e-5, momentum=0.9)

model.compile(loss=tf.keras.losses.Huber(),
              optimizer=optimizer,
              metrics=['mae'])

rnn_forecast = model_forecast(model, series[..., np.newaxis], window_size)
rnn_forecast = rnn_forecast[split_time - window_size:-1, -1, 0]
```

```
model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv1d (Conv1D)	(None, None, 60)	360
lstm (LSTM)	(None, None, 60)	29040
lstm_1 (LSTM)	(None, None, 60)	29040
dense (Dense)	(None, None, 30)	1830
dense_1 (Dense)	(None, None, 10)	310
dense_2 (Dense)	(None, None, 1)	11
lambda (Lambda)	(None, None, 1)	0

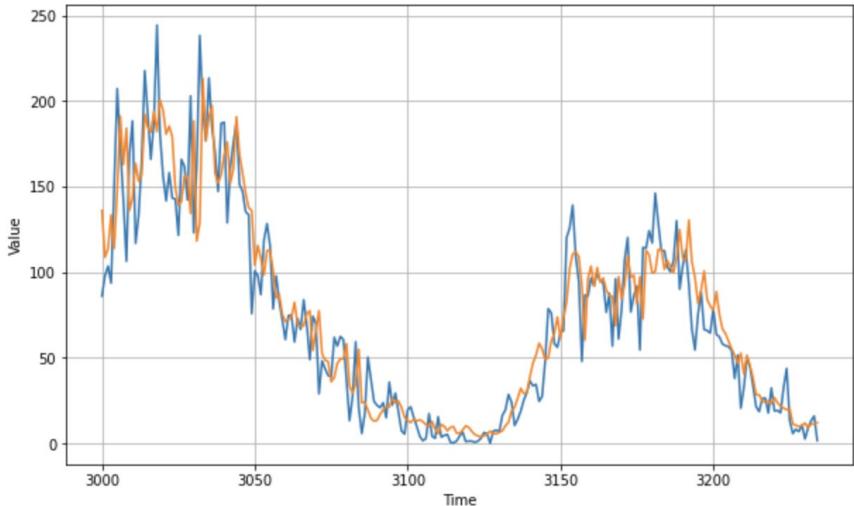
Total params: 60,591

Trainable params: 60,591

Non-trainable params: 0

```
history = model.fit(train_set, epochs=500)
```

Real-world time series data

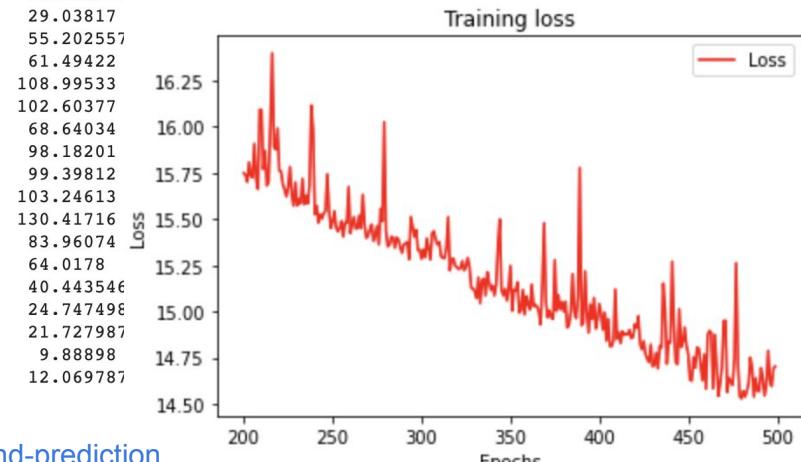
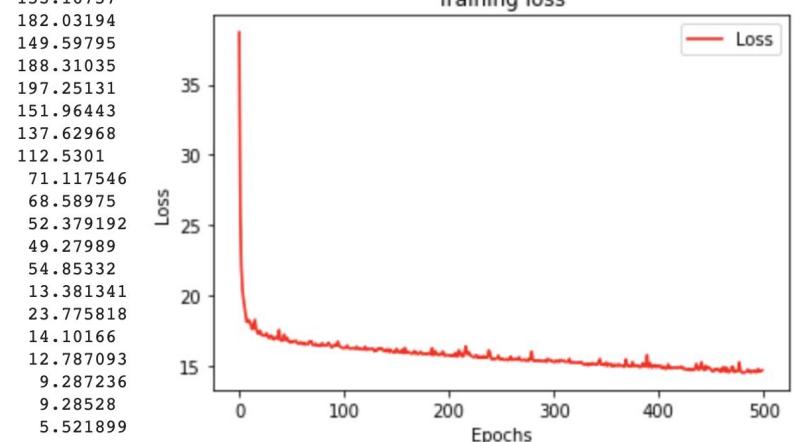


```
tf.keras.metrics.mean_absolute_error(x_valid, rnn_forecast).numpy()
```

15.171835

```
print(rnn_forecast)
```

```
[136.03911 108.64892 113.79592 133.36797 113.831924 148.42674  
191.09799 162.74585 184.07158 135.99425 143.39186 163.61166  
153.16757  
182.03194  
149.59795  
188.31035  
197.25131  
151.96443  
137.62968  
112.5301  
71.117546  
68.58975  
52.379192  
49.27989  
54.85332  
13.381341  
23.775818  
14.10166  
12.787093  
9.287236  
9.28528  
5.521899  
10.305211  
29.03817  
55.202557  
61.49422  
108.99533  
102.60377  
68.64034  
98.18201  
99.39812  
103.24613  
130.41716  
83.96074  
64.0178  
40.443546  
24.747498  
21.727987  
9.88898  
12.069787
```



Source: <https://www.coursera.org/learn/tensorflow-sequences-time-series-and-prediction>

Real-world time series data

Screenshot of a GitHub repository page titled "Datasets". The repository has 26 stars and 402 forks. It contains machine learning datasets used in tutorials on MachineLearningMastery.com, with a link to the website (<http://MachineLearningMastery.com>). The repository includes categories for Binary Classification Datasets, Multiclass Classification Datasets, and Regression Datasets.

- Binary Classification Datasets**
 - Breast Cancer (Wisconsin) ([breast-cancer-wisconsin.csv](#))
 - Breast Cancer (Yugoslavia) ([breast-cancer.csv](#))
 - Breast Cancer (Haberman's) ([haberman.csv](#))
 - Bank Note Authentication ([banknote_authentication.csv](#))
 - Horse Colic ([horse-colic.csv](#))
 - Ionosphere ([ionosphere.csv](#))
 - Pima Indians Diabetes ([pima-indians-diabetes.csv](#))
 - Sonar Returns ([sonar.csv](#))
 - German Credit ([german.csv](#))
 - Credit Card Fraud ([creditcard.csv.zip](#))
 - Adult Income ([adult-all.csv](#))
 - Mammography ([mammography.csv](#))
 - Oil Spill ([oil-spill.csv](#))
 - Phoneme ([phoneme.csv](#))
- Multiclass Classification Datasets**
 - Glass Identification ([glass.csv](#))
 - Iris Flower Species ([iris.csv](#))
 - Wheat Seeds ([wheat-seeds.csv](#))
 - Wine ([wine.csv](#))
 - Ecoli ([ecoli.csv](#))
 - Thyroid Gland ([new-thyroid.csv](#))
- Regression Datasets**
 - Boston Housing ([housing.csv](#))
 - Auto Insurance Total Claims ([auto-insurance.csv](#))
 - Auto Imports Prices ([auto_imports.csv](#))
 - Abalone Age ([abalone.csv](#))
 - Wine Quality Red ([winequality-red.csv](#))
 - Wine Quality White ([winequality-white.csv](#))

Source: <https://github.com/jbrownlee/Datasets>

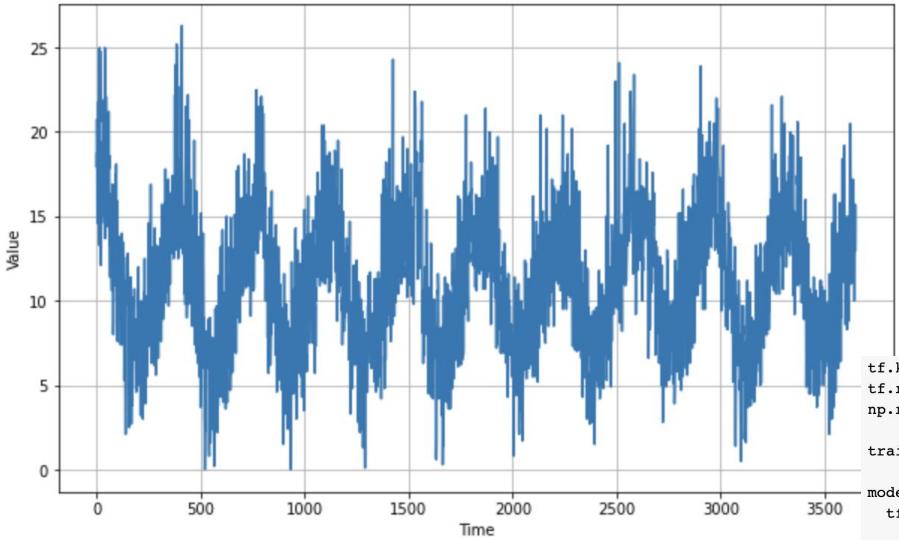
Univariate Time Series Datasets

- Daily Minimum Temperatures in Melbourne ([daily-min-temperatures.csv](#))
- Daily Maximum Temperatures in Melbourne ([daily-max-temperatures.csv](#))
- Daily Female Births in California ([daily-total-female-births.csv](#))
- Monthly International Airline Passengers ([monthly-airline-passengers.csv](#))
- Monthly Armed Robberies in Boston ([monthly-robberies.csv](#))
- Monthly Sunspots ([monthly-sunspots.csv](#))
- Monthly Champagne Sales ([monthly_champagne_sales.csv](#))
- Monthly Shampoo Sales ([monthly-shampoo-sales.csv](#))
- Monthly Car Sales ([monthly-car-sales.csv](#))
- Monthly Mean Temperatures in Nottingham Castle ([monthly-mean-temp.csv](#))
- Monthly Specialty Writing Paper Sales ([monthly-writing-paper-sales.csv](#))
- Yearly Water Usage in Baltimore ([yearly-water-usage.csv](#))

Multivariate Time Series Datasets

- Hourly Pollution Levels in Beijing
- Minutely Individual Household Electric Power Consumption
- Human Activity Recognition Using Smartphones
- Indoor Movement Prediction
- Yearly Longley Economic Employment

Real-world time series data



```
split_time = 2500
time_train = time[:split_time]
x_train = series[:split_time]
time_valid = time[split_time:]
x_valid = series[split_time:]

window_size = 30
batch_size = 32
shuffle_buffer_size = 1000
```

Branch: master Datasets / daily-min-temperatures.csv Find file Copy path

Jason Brownlee Added more time series datasets. 5224989 on Mar 10, 2018 0 contributors

3651 lines (3651 sloc) | 66.3 KB Raw Blame History

Search this file...

	Date	Temp
1	1981-01-01	20.7
2	1981-01-02	17.9
3	1981-01-03	18.8
4	1981-01-04	14.6
5	1981-01-05	15.8
6	1981-01-06	15.8
7	1981-01-07	15.8

```
tf.keras.backend.clear_session()
tf.random.set_seed(51)
np.random.seed(51)

train_set = windowed_dataset(x_train, window_size=60, batch_size=100, shuffle_buffer=shuffle_buffer_size)

model = tf.keras.models.Sequential([
    tf.keras.layers.Conv1D(filters=32, kernel_size=5,
                          strides=1, padding="causal",
                          activation="relu",
                          input_shape=[None, 1]),
    tf.keras.layers.LSTM(64, return_sequences=True),
    tf.keras.layers.LSTM(64, return_sequences=True),
    tf.keras.layers.Dense(30, activation="relu"),
    tf.keras.layers.Dense(10, activation="relu"),
    tf.keras.layers.Dense(1),
    tf.keras.layers.Lambda(lambda x: x * 400)
])

optimizer = tf.keras.optimizers.SGD(lr=1e-5, momentum=0.9)

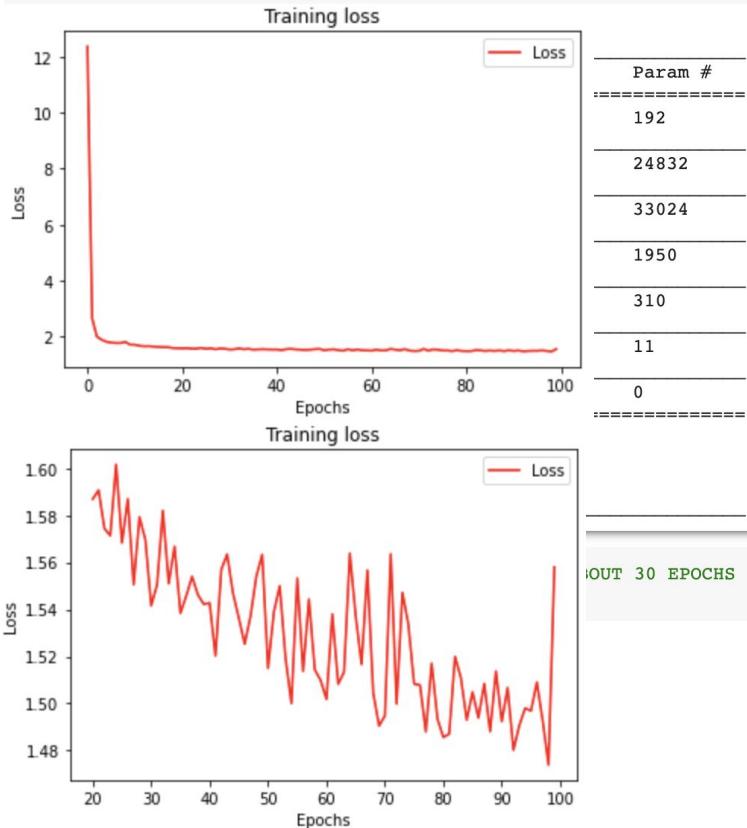
model.compile(loss=tf.keras.losses.Huber(),
              optimizer=optimizer,
              metrics=["mae"])

plt.semilogx(history.history["lr"], history.history["loss"])
plt.axis([1e-8, 1e-4, 0, 60])
(1e-08, 0.0001, 0.0, 60.0)
```

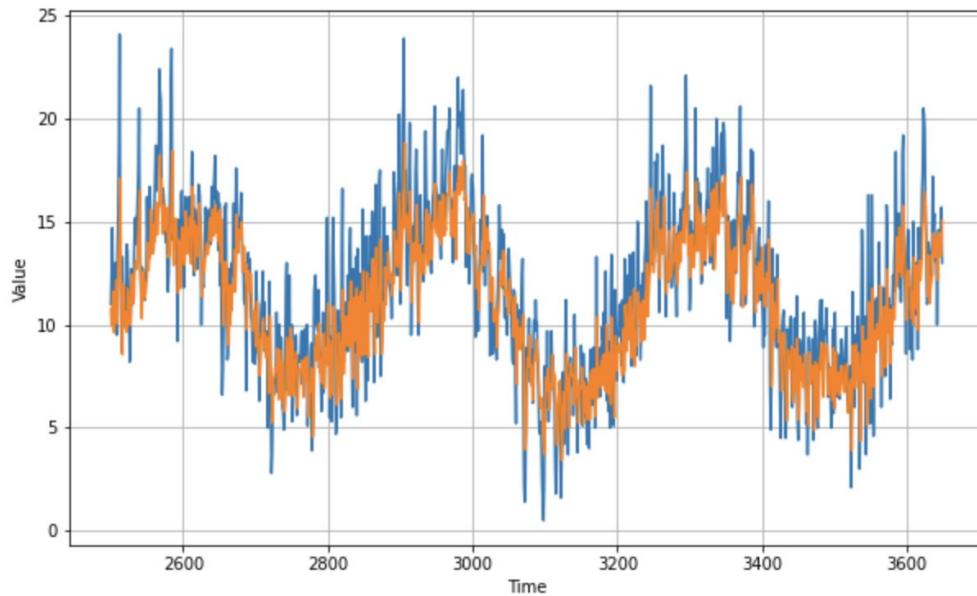
A line plot showing the learning rate and loss over time. The x-axis is logarithmic, ranging from 10^{-8} to 10^{-4} . The y-axis ranges from 0 to 60. The plot shows a sharp decrease in loss as the learning rate decreases, followed by a small peak in loss at a very low learning rate.

Real-world time series data

```
model.summary()
```



```
rnn_forecast = model_forecast(model, series[...], np.newaxis), window_size)  
rnn_forecast = rnn_forecast[split_time - window_size:-1, -1, 0]
```



```
tf.keras.metrics.mean_absolute_error(x_valid, rnn_forecast).numpy()
```

```
# EXPECTED OUTPUT MAE < 2 -- I GOT 1.789626
```

```
1.8799866
```

Course 4: Sequences, Time Series and Prediction

Week 1: Sequences and Prediction

Week 2: Deep Neural Networks for Time Series

Week 3: Recurrent Neural Networks for Time Series

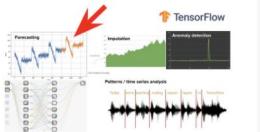
Week 4: Real-world time series data

Let's continue our Time Series adventure 😊

FRI, JUN 19, 7:30 PM EDT

TensorFlow in Practice - C4 Week 1,2 - Forecasting and DNN for Time...

Online event



Join us for our 9th adventure in Deep Learning! Just bring your curiosity and get ready to meet our growing community 😊 We are taking Course 4 of TensorFlow in Practice Specialization available at:...



39 attendees



31 attendees

Attend

FRI, JUL 3, 7:30 PM EDT

A chat with Laurence Moroney, AI Lead at Google

Online event



Join us for a fun conversation with Laurence Moroney, AI Lead at Google (<https://www.linkedin.com/in/laurence-moroney>) and developer of our TensorFlow in Practice and TensorFlow: Data and Deployment Specializations 🎉 We plan to...



9 attendees

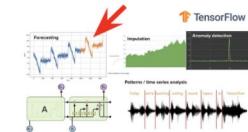
1

Attend

FRI, JUN 26, 7:30 PM EDT

TensorFlow in Practice - C4 Week 3,4 - Forecasting & RNN, Conv1D f...

Online event



Join us for our 10th adventure in Deep Learning! Just bring your curiosity and get ready to meet our growing community 😊 We are taking our last Course 4 of TensorFlow in Practice Specialization available at:...

Attend

FRI, JUL 10, 7:30 PM EDT

How to prepare for and pass the TensorFlow Developer Certificate 🎉

Online event



14 attendees

Attend

Check out these resources

A DISCIPLINED APPROACH TO NEURAL NETWORK HYPER-PARAMETERS: PART 1 – LEARNING RATE, BATCH SIZE, MOMENTUM, AND WEIGHT DECAY

Leslie N. Smith

US Naval Research Laboratory
Washington, DC, USA
leslie.smith@nrl.navy.mil

 Interested

 Add to Calendar

 Share

Tue Jun 16, 9:00 AM - 12:00 PM EDT (3 Hours)

Training Session: Beyond AutoPilot

How to level up your modeling after running Autopilot. Learning about new techniques, algorithms, and features DataRobot makes available to improve your models.



ABOUT M COMPETITIONS M CONFERENCES M PUBLICATIONS M COURSES CONTACT

Home / Dataset

The Dataset

The M5 consists of 10,000 time series of Yearly, Quarterly, Monthly and Other (Weekly, Daily and Hourly) data.

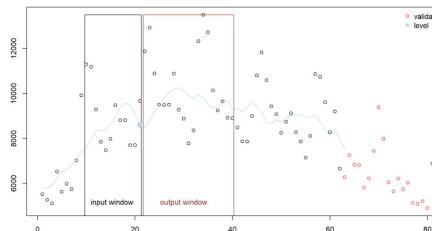
The minimum number of observations is 13 for yearly, 16 for quarterly, 42 for monthly, 80 for weekly, 93 for daily and 700 for hourly series.

The 10,000 time series of the dataset come mainly from the Economic, Finance, Demographics and Industry areas, while also including data from Tourism, Trade, Labor and Wage, Real Estate, Transportation, Natural Resources and the Environment.

The M5 Competition series, as those of the M1 and M3, aim at representing the real world as much as possible. The series were selected randomly from a database of 900,000 ones on December 28, 2017. Professor Makridakis chose the seed number for generating the random sample that determined the M4 Competition data. Some pre-defined filters were applied beforehand to achieve some desired characteristics, such as the length of the series, the percentage of Yearly, Quarterly, Monthly, Weekly, Daily, and Hourly data, as well as their type (Micro, Macro, Finance, Industry, Demographic, Other).

M4 Forecasting Competition: Introducing a New Hybrid ES-RNN Model

Slawek Smyl, Jai Ranganathan, and Andrea Pasqua



June 25, 2018

Sign up for Uber Engineering updates:

Your email address

Subscribe

Popular Articles



Meet Michelangelo: Uber's Machine Learning Platform
September 5, 2017



Uber's Big Data Platform: 100+ Petabytes with Minute Latency
October 17, 2018

Source: <https://arxiv.org/pdf/1803.09820.pdf>

<https://aiworldwide.datarobot.com/agenda/session/259782>

<https://mofc.unic.ac.cy/the-dataset/>

<https://eng.uber.com/m4-forecasting-competition/>

Check out these resources



Get more from Colab

UPGRADE NOW

\$9.99/month

Recurring billing • Cancel anytime

Restrictions apply, learn more here



Faster GPUs

Priority access to faster GPUs and TPUs means you spend less time waiting while code is running. [Learn more](#)



Longer runtimes

Longer running notebooks and fewer idle timeouts mean you disconnect less often. [Learn more](#)



More memory

More RAM means better performance, and less running out of memory. [Learn more](#)

[See what Colab Pro benefits would look like](#)



Install

Learn

API

Resources

Community Why TensorFlow

Search

TensorFlow Core

Overview

Tutorials

Guide

TF 1

TensorFlow tutorials

Quickstart for beginners
Quickstart for experts

BEGINNER

ML basics with Keras

Load and preprocess data

Estimator

ADVANCED

Customization

Distributed training

Images

Text

Structured data

Text

Word embeddings

Text classification with an RNN

Text generation with an RNN

Neural machine translation with attention

Image captioning

Transformer model for language understanding

The TensorFlow tutorials are written as Jupyter notebooks and run directly in Google Colab—a hosted notebook environment that requires no setup. Click the [Run in Google Colab](#) button.

For beginners

The best place to start is with the user-friendly Keras sequential API. Build models by plugging together building blocks. After these tutorials, read the [Keras guide](#).

Beginner quickstart

This "Hello, World!" notebook shows the Keras Sequential API and `model.fit`.

Keras basics

This notebook collection demonstrates basic machine learning tasks using Keras.

Load data

These tutorials use `tf.data` to load various data formats and build input pipelines.

For experts

The Keras functional and subclassing APIs provide a define-by-run interface for customization and advanced research. Build your model, then write the forward and backward pass. Create custom layers, activations, and training loops.

Advanced quickstart

This "Hello, World!" notebook uses the Keras subclassing API and a custom training loop.

Customization

This notebook collection shows how to build custom layers and training loops in TensorFlow.

Distributed training

Distribute your model training across multiple GPUs, multiple machines or TPUs.

The Advanced section has many instructive notebook examples, including [Neural machine translation](#), [Transformers](#), and [CycleGAN](#).

Source:

<https://colab.research.google.com/signup>

<https://www.tensorflow.org/tutorials/>

Check out these AI for Healthcare resources



Enroll for Free
Starts May 22

About How It Works Courses Instructors Enrollment Options FAQ

There are 3 Courses in this Specialization

COURSE

AI for Medical Diagnosis

1

★★★★★ 4.6 410 ratings • 102 reviews

AI is transforming the practice of medicine. It's helping doctors diagnose patients more accurately, make predictions about patients' future health, and recommend better treatments. As an AI practitioner, you have the opportunity to join in this transformation of modern medicine. If you're already familiar with some of the math and coding behind AI

[SHOW ALL](#)

COURSE

AI for Medical Prognosis

2

★★★★★ 4.6 113 ratings • 27 reviews

AI is transforming the practice of medicine. It's helping doctors diagnose patients more accurately, make predictions about patients' future health, and recommend better treatments. This Specialization will give you practical experience in applying machine learning to concrete problems in medicine.

[SHOW ALL](#)

COURSE

AI For Medical Treatment

3

AI is transforming the practice of medicine. It's helping doctors diagnose patients more accurately, make predictions about patients' future health, and recommend better treatments. This Specialization will give you practical experience in applying machine learning to concrete problems in medicine.

[SHOW ALL](#)

AI for Healthcare

Learn to build, evaluate, and integrate predictive models that have the power to transform patient outcomes. Begin by classifying and segmenting 2D and 3D medical images to augment diagnosis and then move on to modeling patient outcomes with electronic health records to optimize clinical trial testing decisions. Finally, build an algorithm that uses data collected from wearable devices to estimate the wearer's pulse rate in the presence of motion.

PREREQUISITE KNOWLEDGE

Intermediate Python, and Experience with Machine Learning See detailed requirements.

[HIDE DETAILS](#)



Applying AI to 2D Medical Imaging Data

Learn the fundamental skills needed to work with 2D medical imaging data and how to use AI to derive clinically-relevant insights from data gathered via different types of 2D medical imaging such as x-ray, mammography, and digital pathology. Extract 2D images from DICOM files and apply the appropriate tools to perform exploratory data analysis on them. Build different AI models for different clinical scenarios that involve 2D images and learn how to position AI tools for regulatory approval.

PNEUMONIA DETECTION FROM CHEST X-RAYS

Applying AI to 3D Medical Imaging Data

Learn the fundamental skills needed to work with 3D medical imaging datasets and frame insights derived from the data in a clinically relevant context. Understand how these images are acquired, stored in clinical archives, and subsequently read and analyzed. Discover how clinicians use 3D medical images in practice and where AI holds most potential in their work with these images. Design and apply machine learning algorithms to solve the challenging problems in 3D medical imaging and how to integrate the algorithms into the clinical workflow.

HIPPOCAMPUS VOLUME QUANTIFICATION FOR ALZHEIMER'S PROGRESSION

Applying AI to EHR Data

Learn the fundamental skills to work with EHR data and build and evaluate compliant, interpretable models. You will cover EHR data privacy and security standards, how to analyze EHR data and avoid common challenges, and cover key industry code sets. By the end of the course, you will have the skills to analyze an EHR dataset, transform it to the right level, build powerful features with TensorFlow, and model the uncertainty and bias with TensorFlow Probability and Aequitas.

PATIENT SELECTION FOR DIABETES DRUG TESTING

Applying AI to Wearable Device Data

Learn how to build algorithms that process the data collected by wearable devices and surface insights about the wearer's health. Cover the sensors and signal processing foundation that are critical for success in this domain, including IMU, PPG, and ECG that are common to most wearable devices, and learn how to build three algorithms from real-world sensor data.

MOTION COMPENSATED PULSE RATE ESTIMATION

Source:

<https://www.coursera.org/specializations/ai-for-medicine>

<https://www.udacity.com/course/ai-for-healthcare-nanodegree--nd320>

Check out this certification and books

TensorFlow Core

Introducing the TensorFlow Developer Certificate!

March 12, 2020



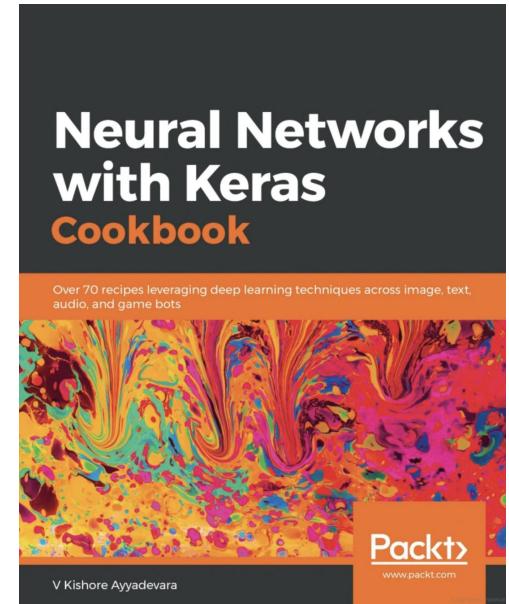
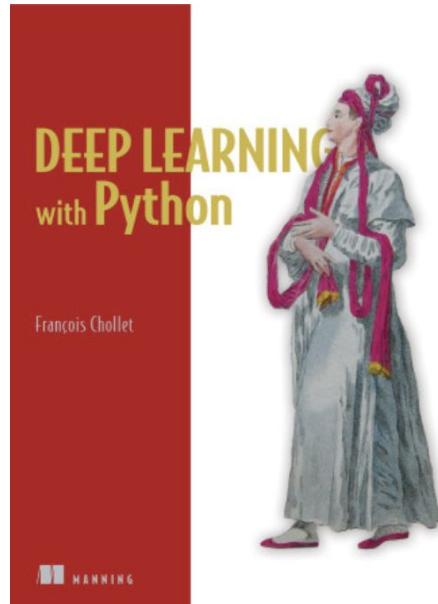
Posted by Alina Shinkarsky, on behalf of the TensorFlow Team

In the AI world today, more and more companies are looking to hire machine learning talent, and simultaneously, an increasing number of students and developers are looking for ways to gain and showcase their ML knowledge with formal recognition. In addition to the courses and learning resources available online, we want to help developers showcase their ML proficiency and help companies hire ML developers to solve challenging problems.



Source:

- <https://blog.tensorflow.org/2020/03/introducing-tensorflow-developer-certificate.html?m=1>
- <https://www.manning.com/books/deep-learning-with-python>
- https://books.google.com/books?id=5quLDwAAQBAJ&printsec=frontcover&source=gbs_qe_summary_r&cad=0#v=onepage&q&f=false
- <https://github.com/PacktPublishing/Neural-Networks-with-Keras-Cookbook/tree/master/Chapter11>

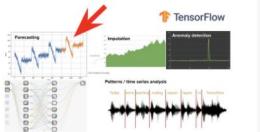


Let's continue our Time Series adventure 😊

FRI, JUN 19, 7:30 PM EDT

TensorFlow in Practice - C4 Week 1,2 - Forecasting and DNN for Time...

Online event



Join us for our 9th adventure in Deep Learning! Just bring your curiosity and get ready to meet our growing community 😊 We are taking Course 4 of TensorFlow in Practice Specialization available at:...



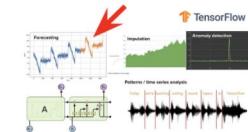
39 attendees



FRI, JUN 26, 7:30 PM EDT

TensorFlow in Practice - C4 Week 3,4 - Forecasting & RNN, Conv1D f...

Online event



Join us for our 10th adventure in Deep Learning! Just bring your curiosity and get ready to meet our growing community 😊 We are taking our last Course 4 of TensorFlow in Practice Specialization available at:...

[Attend](#)

FRI, JUL 3, 7:30 PM EDT

A chat with Laurence Moroney, AI Lead at Google

Online event



Join us for a fun conversation with Laurence Moroney, AI Lead at Google (<https://www.linkedin.com/in/laurence-moroney>) and developer of our TensorFlow in Practice and TensorFlow: Data and Deployment Specializations 🎉 We plan to...



1

[Attend](#)

FRI, JUL 10, 7:30 PM EDT

How to prepare for and pass the TensorFlow Developer Certificate 🎉

Online event

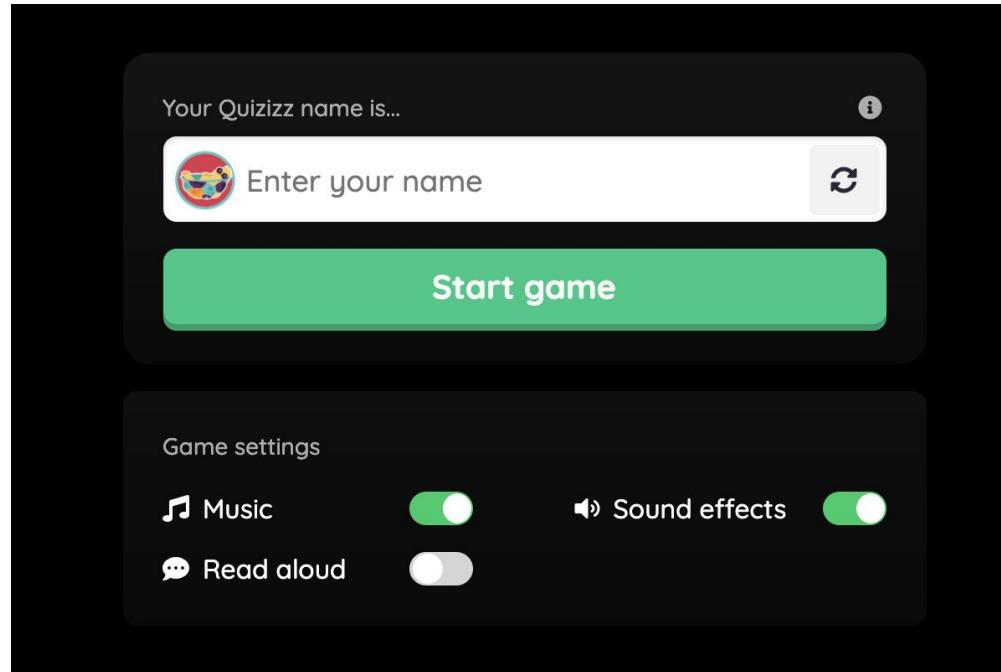
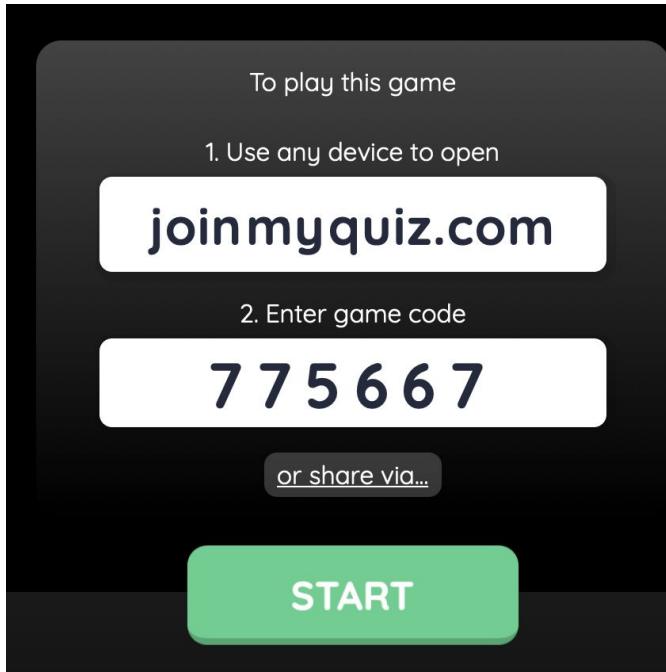


Join us for our 11th adventure in Deep Learning! Just bring your curiosity and get ready to meet our growing community 😊 Join Zoom Meeting: <https://us02web.zoom.us/j/84402592502?...>



[Attend](#)

Time for a fun game 😊🎉



Practice here or use Flashcards:

<https://quizizz.com/join/quiz/5eefd13d3db4b8001bb6d72b/start?from=soloLinkShare&referrer=5d921444d0fa99001a135336>

Time for a fun game



Practice here or use Flashcards:

<https://quizizz.com/join/quiz/5eefd13d3db4b8001bb6d72b/start?from=soloLinkShare&referrer=5d921444d0fa99001a135336>

The image shows two side-by-side screenshots. On the left is the Quizizz game dashboard for a game with code 432 236. It displays a list of 22 players with their names, profile icons, scores, and progress bars. A video feed of a person is shown in the top right corner. On the right is a slide from a presentation titled "Deep-Learning-Adventures-Chapter-1-Presentation-1 - Google Slides". The slide has a purple header with the question "How do you import Tensorflow in your Python code". Below the question are four numbered options: 1. import tf from ai, 2. import tensorflow from google, 3. expot tensorflow as tf, and 4. import tensorflow as tf.

Rank	Name	Score
6	Anil	8630
6	rek	8630
7	Charles Stockman	7450
8	Martin	7290
9	Angel	7230
10	Chuba	7050
11	Peter	6880
12	Dean	6840
13	Melissa	6505
14	Sanjay	5710
15	Hasson	4750
16	AH	4470
16	Kottie	4470
17	v	4010
18	KL	3910
19	Tony	3240
20	SS	2130
21	George	2100
22	Thomas	0

Questions

Discussion

2 Deep Learning representations

For representations:

- nodes represent inputs, activations or outputs
- edges represent weights or biases

Here are several examples of Standard deep learning representations

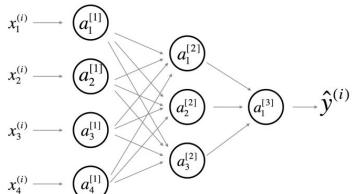


Figure 1: Comprehensive Network: representation commonly used for Neural Networks. For better aesthetic, we omitted the details on the parameters ($w_{ij}^{[l]}$ and $b_i^{[l]}$ etc...) that should appear on the edges

Standard notations for Deep Learning

This document has the purpose of discussing a new standard for deep learning mathematical notations.

1 Neural Networks Notations.

General comments:

- superscript (i) will denote the i^{th} training example while superscript [l] will denote the l^{th} layer

Sizes:

$\cdot m$: number of examples in the dataset

$\cdot n_x$: input size

$\cdot n_y$: output size (or number of classes)

$\cdot n_h^{[l]}$: number of hidden units of the l^{th} layer

In a for loop, it is possible to denote $n_x = n_h^{[0]}$ and $n_y = n_h^{[\text{number of layers} + 1]}$.

$\cdot L$: number of layers in the network.

Objects:

$\cdot X \in \mathbb{R}^{n_x \times m}$ is the input matrix

$\cdot x^{(i)} \in \mathbb{R}^{n_x}$ is the i^{th} example represented as a column vector

$\cdot Y \in \mathbb{R}^{n_y \times m}$ is the label matrix

$\cdot y^{(i)} \in \mathbb{R}^{n_y}$ is the output label for the i^{th} example

$\cdot W^{[l]} \in \mathbb{R}^{\text{number of units in next layer} \times \text{number of units in the previous layer}}$ is the weight matrix, superscript [l] indicates the layer

$\cdot b^{[l]} \in \mathbb{R}^{\text{number of units in next layer}}$ is the bias vector in the l^{th} layer

$\cdot \hat{y} \in \mathbb{R}^{n_y}$ is the predicted output vector. It can also be denoted $a^{[L]}$ where L is the number of layers in the network.

Common forward propagation equation examples:

$a = g^{[l]}(W_x x^{(i)} + b_1) = g^{[l]}(z_1)$ where $g^{[l]}$ denotes the l^{th} layer activation function

$\hat{y}^{(i)} = \text{softmax}(W_h h + b_2)$

· General Activation Formula: $a_j^{[l]} = g^{[l]}(\sum_k w_{jk}^{[l]} a_k^{[l-1]} + b_j^{[l]}) = g^{[l]}(z_j^{[l]})$

· $J(x, W, b, y)$ or $J(\hat{y}, y)$ denote the cost function.

Examples of cost function:

$\cdot J_{CE}(\hat{y}, y) = -\sum_{i=0}^m y^{(i)} \log \hat{y}^{(i)}$

$\cdot J_1(\hat{y}, y) = \sum_{i=0}^m |y^{(i)} - \hat{y}^{(i)}|$

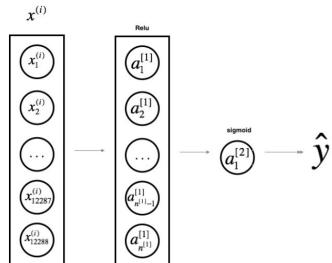


Figure 2: Simplified Network: a simpler representation of a two layer neural network, both are equivalent.